

Making the Computer See

Computer Vision For Everyday Applications

Martin Jul

NDC Oslo, June 5th, 2014

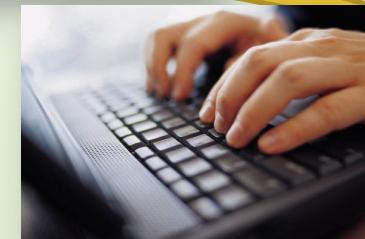


Our Habitual World...

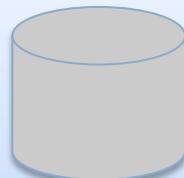
*Unstructured
Data*



Human Data Processing



Structured
Data



Machine Data Processing

*Producer: MF
Wine: Pinot Noir
North Coast
Vintage: 2011*

Fun?



Fun?



<http://en.m.wikipedia.org>

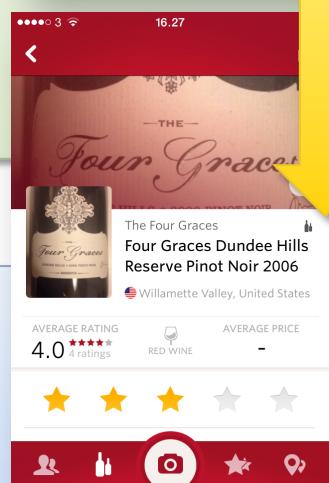
Fun!



A screenshot of a mobile application interface for a wine product. At the top, there is a red header bar with icons for signal strength, battery level, and a back arrow. The main content area shows a large image of the wine bottle. Below the image, the wine's name is displayed in a large, elegant script font: "THE Four Graces". Underneath, the text reads "Four Graces Dundee Hills Reserve Pinot Noir 2006". To the right of the text is a small edit icon (a pencil inside a circle). Further down, the text "Willamette Valley, United States" is shown with a small American flag icon. Below this, there are two sections: "AVERAGE RATING" with a 4.0 rating and 4 ratings, and "AVERAGE PRICE" with a redacted value. A row of five stars follows, with the first three filled yellow and the last two empty grey. At the bottom, there is a red navigation bar with icons for user profile, bottle, camera, star, and location.

“Software is Eating the World”

*Unstructured
Data*

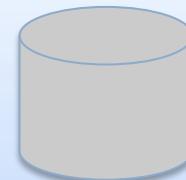


Human Data Processing



Four Graces
Dundas Hills Reserve
Pinot Noir
2006

Structured
Data



Machine Data Processing

Agenda

- Historical Context
- Books and Cars – the 20 mill. kroner problem
 - Recognizing Shapes and Geometry
 - Reading text from images
- Pizza is Here
 - Object Detection
- Wine
 - Computer Vision and Machine Learning

A long time ago in a galaxy far,
far away....

On a Type-reading Optophone.

By E. E. FOURNIER D'ALBE, D.Sc. (Lond. and Birm.).

(Communicated by Sir Oliver Lodge, F.R.S. Received May 2,—Read
May 28, 1914.)

The prod
light is the
“photophon
of substitut
priately ter
Having c

I described
bright ligl
discover s
putting tv
sending th
current b
the action

This disadvantage was
in 1913.† The audible te
of various musical frequ
dots placed in a row, it
5 cm. high by learning
In order to adapt this
letterpress by means of

- (1) The length of th
to about 1.5 mm., the
 - (2) The light had t
 - (3) The sensitivenc
- to be greatly augmen

These modification
which should, with
ordinary books and

* 'Physikalische Z
shown at the Optica
June 25, 1912.

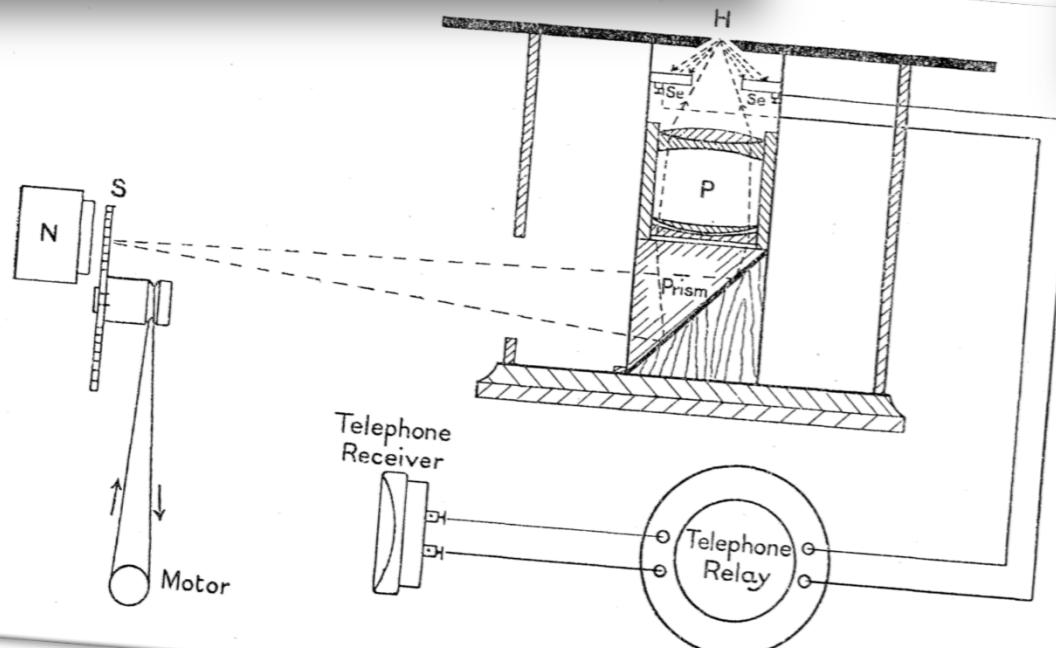
+ 'The Electrician
action at Birmingh
Association.

On a Type-reading Optophone.

By E. E. FOURNIER D'ALBE, D.Sc. (Lond. and Birm.).

(Communicated by Sir Oliver Lodge, F.R.S. Received May 2,—Read
May 28, 1914.)

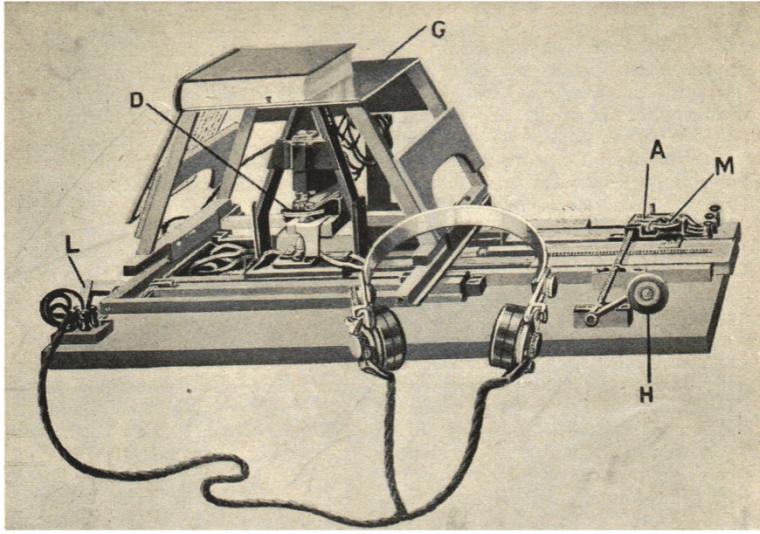
The production of sounds directly or indirectly due to the incidence of light is the general function of instruments of the type of Graham Bell's "photophone." An instrument designed to solve the more special problem of substituting the sense of hearing for the sense of sight is more appropriately termed an "ophone."



1910s



1920s



<https://en.wikipedia.org/wiki/Optophone>

SEVENTY-SIXTH YEAR

SCIENTIFIC AMERICAN

THE WEEKLY JOURNAL OF PRACTICAL INFORMATION

NEW YORK. NOVEMBER 6, 1920

15 CENTS A COPY
20 CENTS IN CANADA

Fig. 1.—Skeleton apparatus showing the principle of the "black-sounding" optophone.

Fig. 2.—The optophone with book not removed, showing the "tracer" mechanism.

Fig. 3.—The optophone complete with conventional book in place ready for reading.

The Type-Reading Optophone
An Instrument Which Enables the Blind To Read Ordinary Type

I

T would be difficult to conceive of a more remarkable electrical device than the type-reading optophone, which enables the blind to read ordinary type. This device, it is understood, is not another new affair as its basic principles are concerned. As far back as August 3rd, 1912, the SCIENTIFIC AMERICAN PRESENT carried a description of the first optophone which in its earliest form was an instrument for converting light into sound in such a manner that the blind could easily see by ear, the invention of Dr. E. G. Fourier d'Alba of Birmingham University. Dr. d'Alba later developed a type-reading instrument by which the blind could read large prints. Recently the instrument has been greatly improved in cooperation with Messrs. Barr and Strong of Glasgow, the well-known instrument makers. In its present form it can be adjusted for any ordinary type, but it is well to direct attention to the fact that the blind person making use of the optophone will have a very fine time in ear in order to grasp the tone combinations emitted by the instrument. The general principle of the optophone is as follows:

In Fig. 1, wherein a siren disk, *D*, is run at about 30 revolutions a second by means of a small magnetoelectric motor, is shown. It contains five circles of square perforations, the innermost circle having 24 perforations, the outermost 42, the other circles being intermediate and corresponding to the relative frequency of certain notes of the diaphragm scale. A line of light in a radial direction is provided by the straight-flame lamp, *L*, and the image of the filament of this lamp is thrown on the print by a system of three lenses on the other side of the selenium tablet *S*. The axis of the concavo-convex lens is slightly tilted out of axis of the other lenses for the purpose which is specifically herein. The general result is that the optical system is to give a line of luminous dots on the print, each dot having a different musical frequency. The condensing tube, which is put in circuit with a battery and a high-resistance telephone receiver, however. Those dots which fall on white paper produce a series of their own musical frequency in the telephone, while those which fall on black are distinguished, so to speak. We thus get what may be called a "white-sounding" optophone, which the black letters are read by the notes omitted from the scale rather than by the notes which remain sounding. All the reading demonstrations hitherto undertaken have been given with a "white-

sounding" optophone, with the most satisfying results.

A modification of this principle, we learn from the SCIENTIFIC AMERICAN MONTHLY [succinctly to the SCIENTIFIC AMERICAN SUPPLEMENT] for October 1912, has been introduced by Messrs. Barr and Strong in connection with Dr. d'Alba. This is in the provision of a second selenium preparation in the form of a cylindrical rod, the total length of which can be seen at *M*. This receives the light reflected from the concave surface of the lens *C*, which produces a real image of the line of dots on *S*, and about as fast the image can be made more or less exact against the spot on *S*, when white paper alone is exposed, a silence, and the effect on *B*, where the filament lamp is brought into focus by the change-bar *C*, which has a friction grip inside the bar on which the "tracer" is pivoted, and can be adjusted for any desired line speed by means of the screw attached to the change bar. A lever attached to the "tracer" enables the operator to reverse the motion or to release the "tracer" from the friction grip, so that it may be quickly brought to the top of the page. The straight-flame lamp is inserted in *L*, where it is held by a spring clip, and whence it can easily be removed for renewal or by a blind operator. The balance screw, inserted at *B*, can also be adjusted for balance by means of the small handle shown.

Fig. 3 shows the apparatus from the top page end, with telephone flexible connections attached, as well as the book-rest. For reading a book. The adapter has flexible connections all of different sizes and fit into different-sized holes in such a manner that they cannot be easily inserted.

The operator of the optophone reads by the tones and combinations of tones emitted by the telephone receiver. The numbers of perforations in the siren disk are in proportion to the notes *G*, *C*, *D'*, *E*, *F*, *G'*, *A*, *B*, *C*, *D*, *E*, *F*, *G*, *H*, *I*, *J*, *K*, *L*, *M*, *N*, *O*, *P*, *Q*, *R*, *S*, *T*, *U*, *V*, *W*, *X*, *Y*, *Z*, etc., do, re, me, so, mi, fa, la, ti, do, re, me, (high) sol, mi, fa, la, ti, do, re, me, (high) sol, etc.

In practice it is found that, with the new apparatus, the various adjustments for the type, length of lines, and book interval are quite easily made by blind persons, and that the instrument, with all its delicate adjustments, can remain in use for a long time without anything getting out of order.

As for speed of reading, we are told that Miss Mary Mason, the blind girl who some time ago now reads habitually at a speed of about 25 words a minute with a "white-sounding" optophone, thus repeats a so-called musical ear is a very necessary qualification.

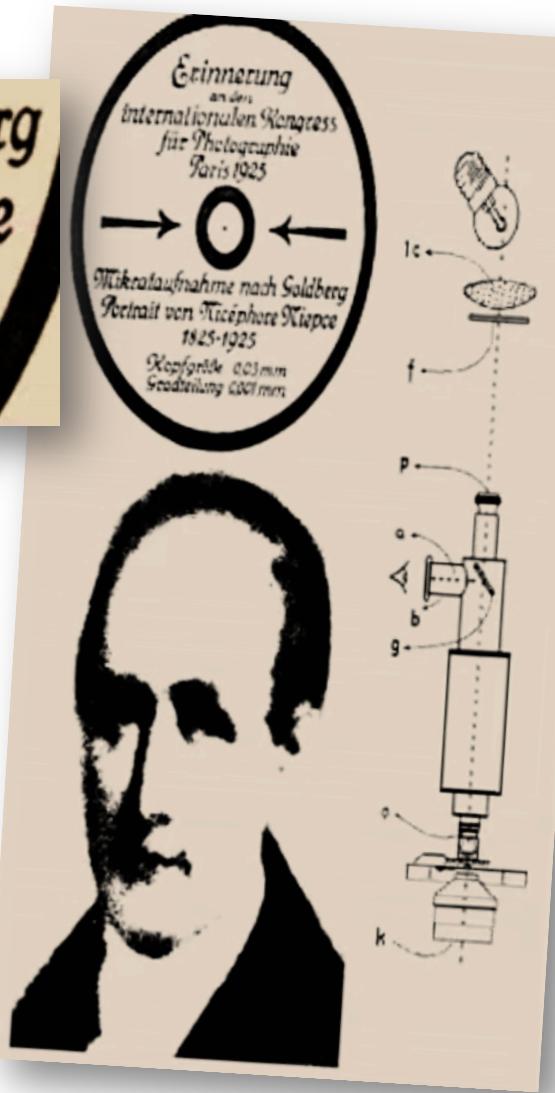
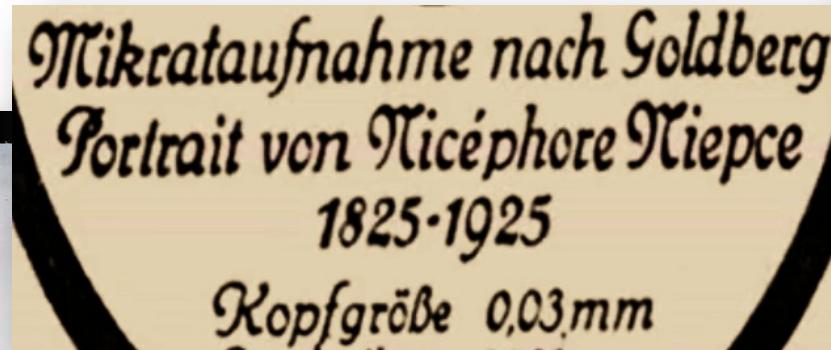
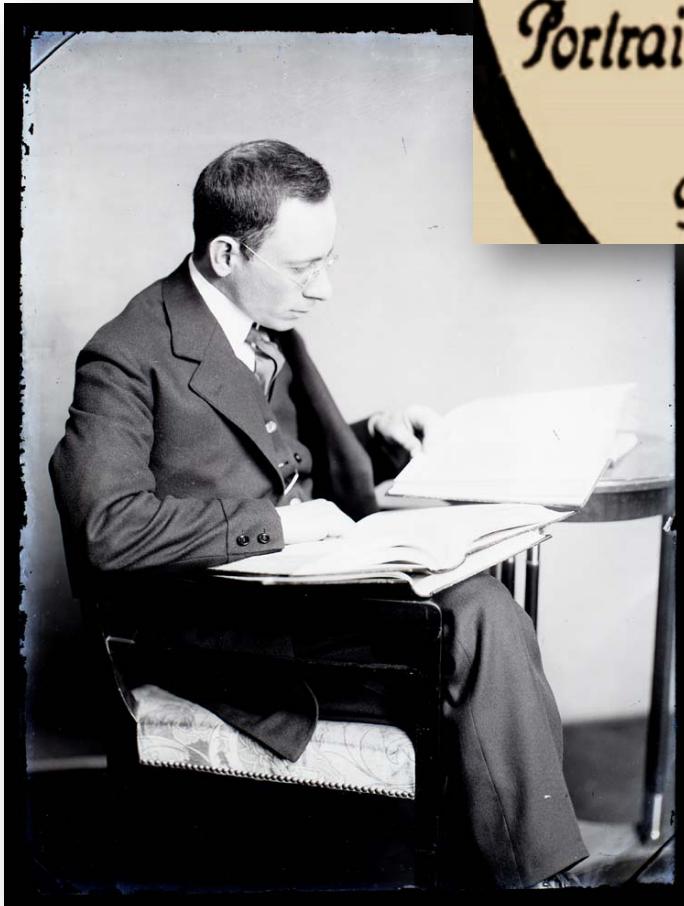
Diagram of the arrangement of the optophone as seen from two sides

Digitized by Google

Original from
UNIVERSITY OF MICHIGAN

Source: <http://babel.hathitrust.org/cgi/pt?id=mdp.39015024546411;view=1up;seq=543>

1920s

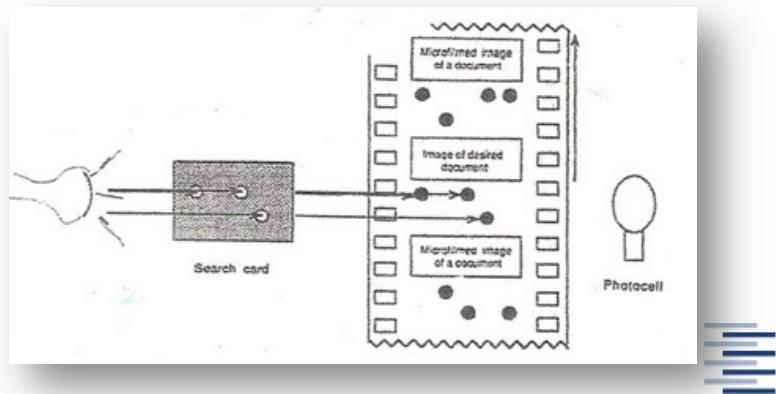


Emanuel Goldberg (1881-1970), photo from 1907,
<http://zeisshistoricasociety.org/emanuel-goldberg/>

1920s – 1930s



Abb. 7-8: Modell der Statistischen Maschine – von oben und der Seite
(Foto M. Turkiewicz 2013, HAW Hamburg)

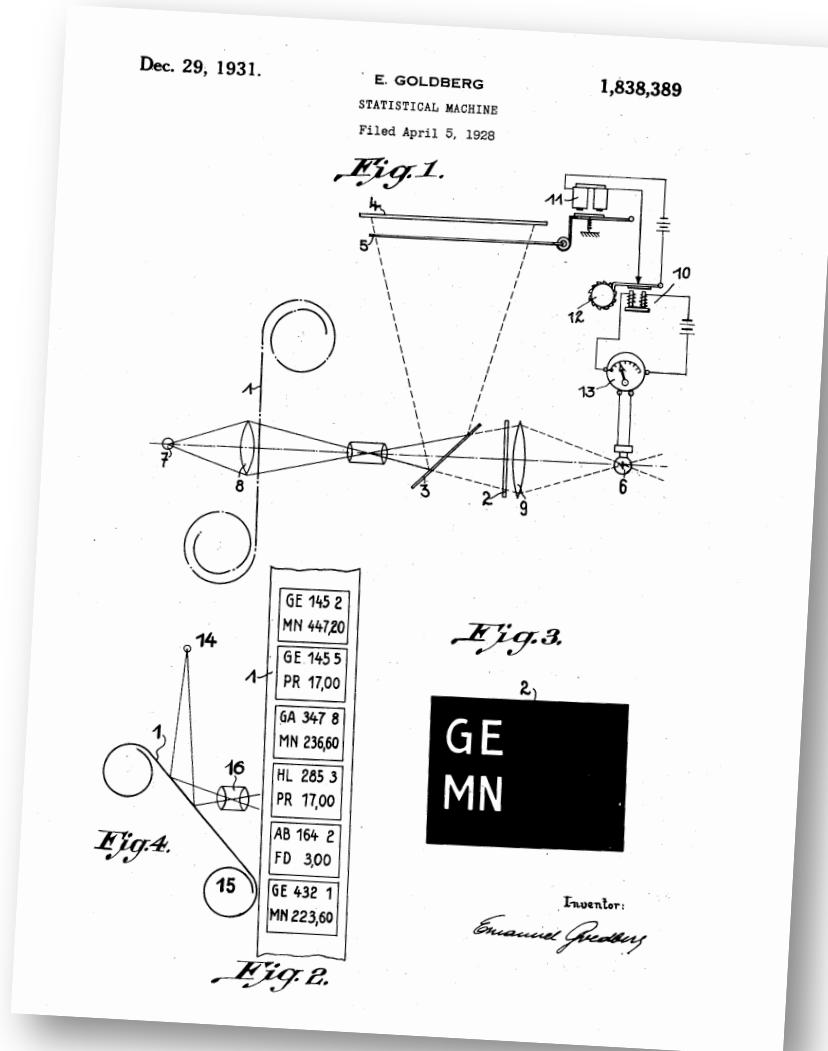


Hochschule für Angewandte Wissenschaften Hamburg
Hamburg University of Applied Sciences

Demonstration model built by HAW Hamburg students:

http://www.searchstudies.org/tl_files/ADSDS/

Projektbericht_Statistische_Maschine_Gurevich_Heuer_Turkiewicz.pdf



Patent info: <http://patft.uspto.gov> patent 1 838 389

1960s

Thule

1910-1960



First day cover



1970s



Kurzweil Reading Machine (1976)

64 KB RAM

¼ MHz CPU

Flatbed CCD Scanner

Software:

- Omnifont OCR
- Text-to-speech

Price in today's money:

195,900 USD (2014 at CPI infl.)

1.16 million NOK

January 1977 – Kurzweil unveils the Reading Machine at Iowa Commission for the Blind

<http://www.iowablindhistory.org/blindhistory/tools-and-technology-optical-character-recognition-scanners>

1980s

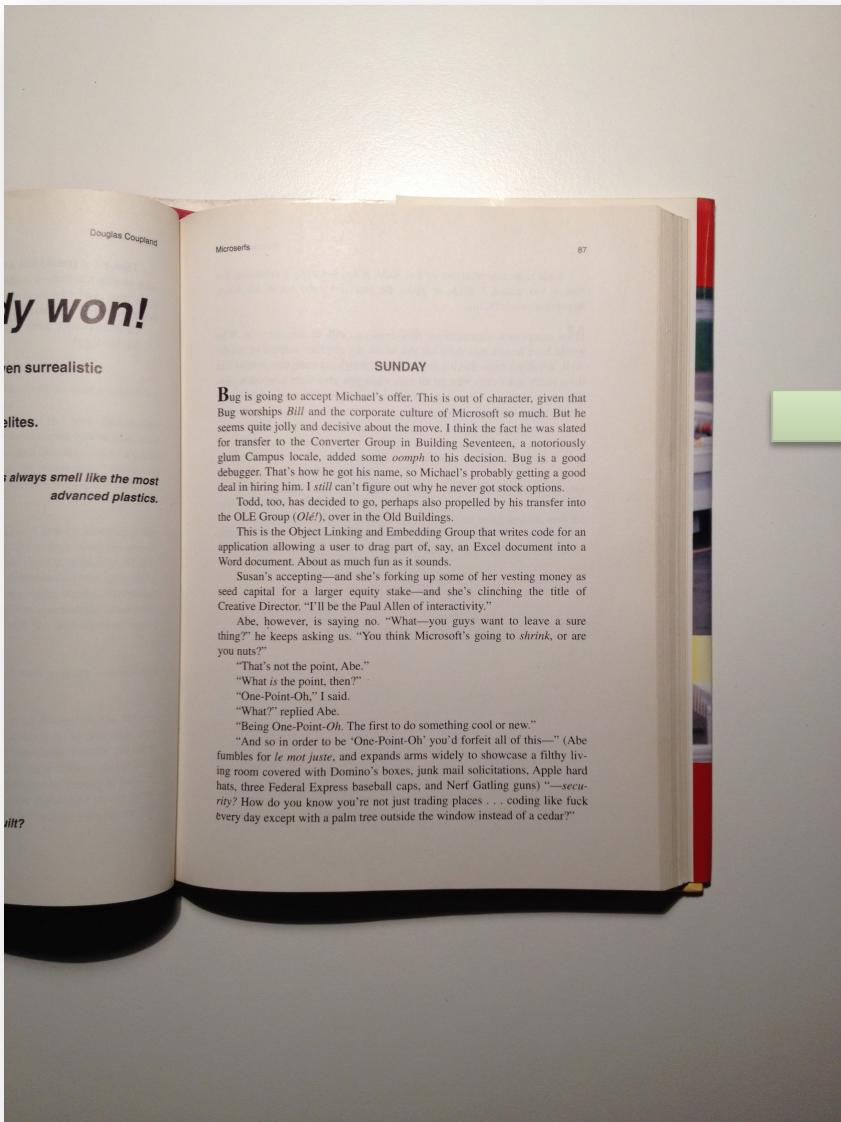


1987 - HP introduces the desktop scanner, ScanJet
http://www.hpmuseum.net/display_item.php?hw=295

OCR

OPTICAL CHARACTER RECOGNITION

Reading Text From Images



SUNDAY

Bug is going to accept Michael's offer. This is out of character, given that Bug worships Bill and the corporate culture of Microsoft so much. But he seems quite jolly and decisive about the move. I think the fact he was slated for transfer to the Converter Group in Building Seventeen, a notoriously glum Campus locale, added some *oomph* to his decision. Bug is a good debugger. That's how he got his name, so Michael's probably getting a good deal in hiring him. I still can't figure out why he never got stock options.

Todd, too, has decided to go, perhaps also propelled by his transfer into the OLE Group (*Ole!*), over in the Old Buildings.

This is the Object Linking and Embedding Group that writes code for an application allowing a user to drag part of, say, an Excel document into a Word document. About as much fun as it sounds.

Susan's accepting—and she's forking up some of her vesting money as seed capital for a larger equity stake—and she's clinching the title of Creative Director. "I'll be the Paul Allen of interactivity."

Abe, however, is saying no. "What—you guys want to leave a sure thing?" he keeps asking us. "You think Microsoft's going to *shrink*, or are you nuts?"

"That's not the point, Abe."

"What is the point then?"

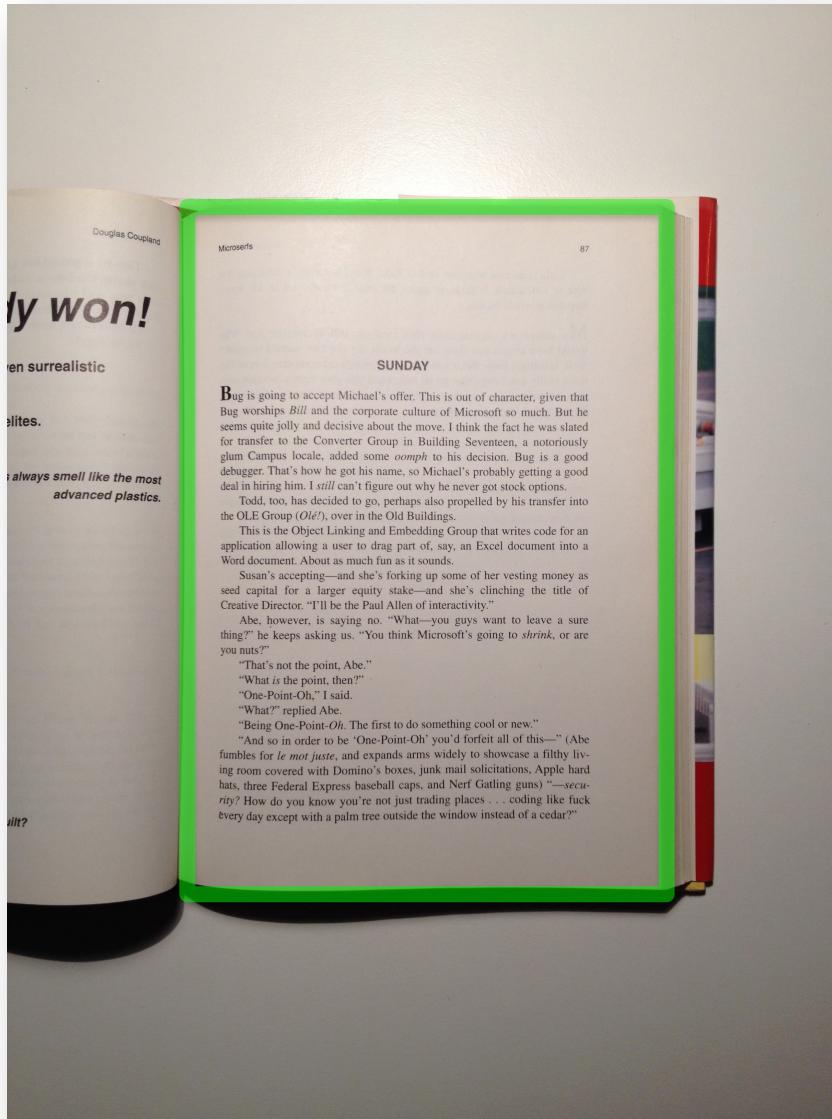
"One-Point-Oh," I said.

"What?" replied Abe.

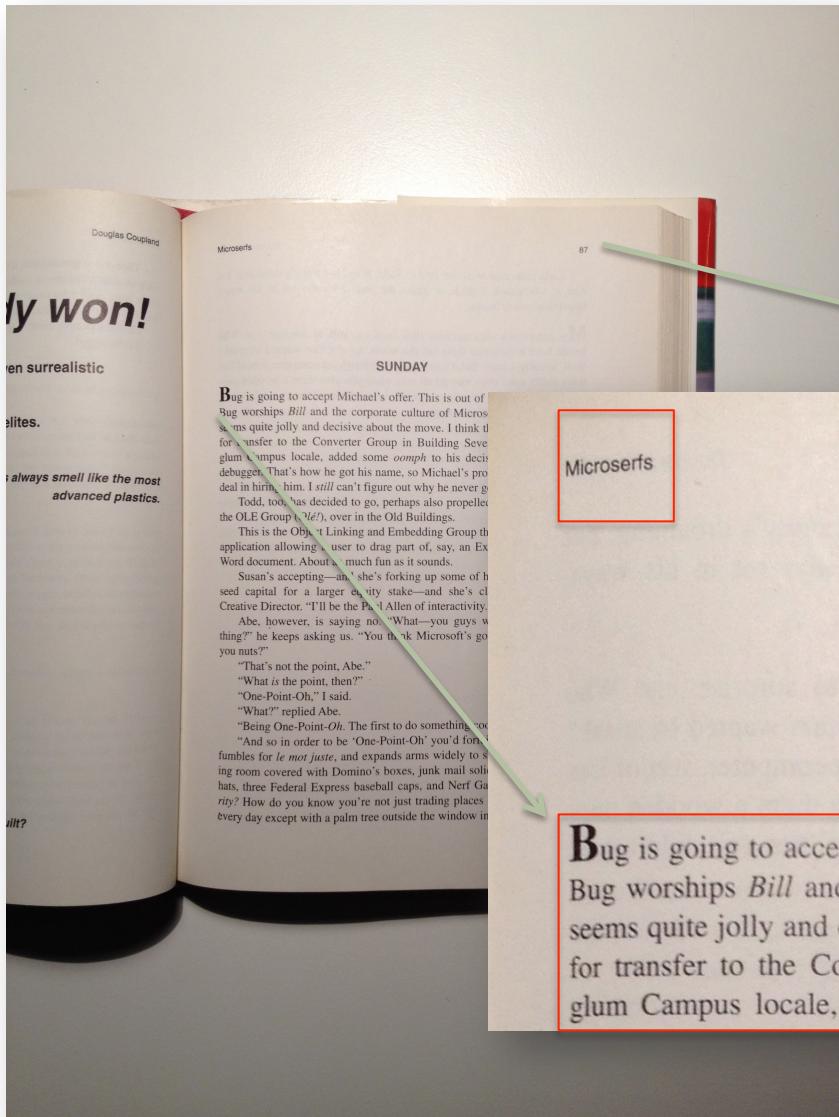
"Being One-Point-Oh. The first to do something cool or new."

"And so in order to be 'One-Point-Oh' you'd forfeit all of this—" (Abe fumbles for *le mot juste*, and expands arms widely to showcase a filthy living room covered with Domino's boxes, junk mail solicitations, Apple hard hats, three Federal Express baseball caps, and Nerf Gatling guns) "—security? How do you know you're not just trading places . . . coding like fuck every day except with a palm tree outside the window instead of a cedar?"

Where is the Page?

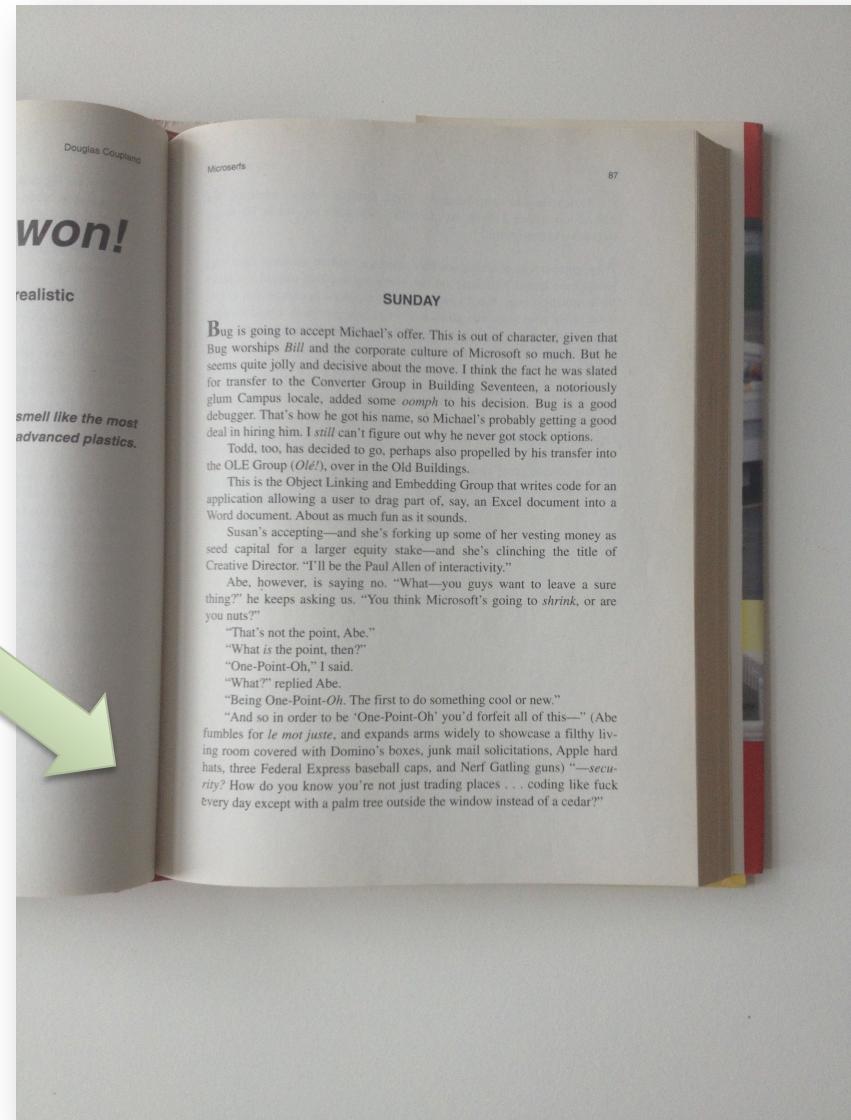
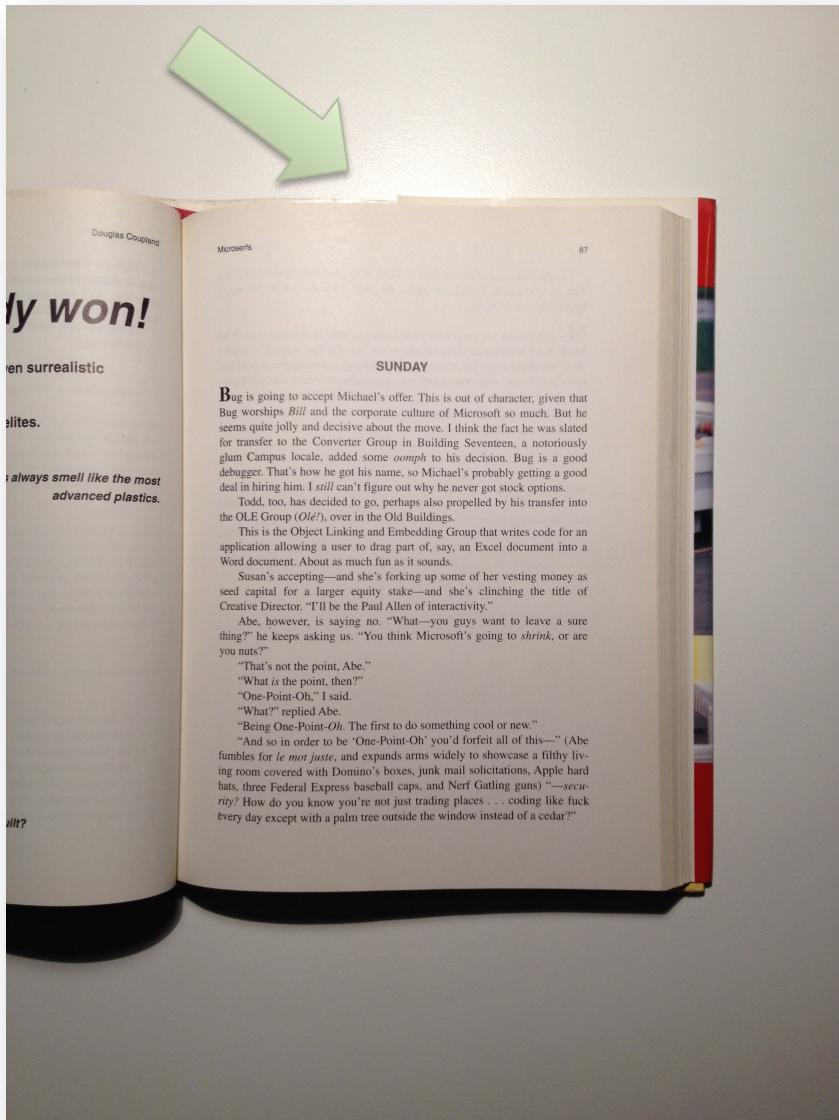


Multiple Texts and Fonts



Bug is going to accept Michael's offer. This is out of character, given that Bug worships Bill and the corporate culture of Microsoft so much. But he seems quite jolly and decisive about the move. I think the fact he was slated for transfer to the Converter Group in Building Seventeen, a notoriously glum Campus locale, added some *oomph* to his decision. Bug is a good deal in him, him. I still can't figure out why he never gets off the bus. Todd, too, has decided to go, perhaps also propelled by the Ole Group (*Idiot*), over in the Old Buildings.

White Isn't



Straight Lines Aren't

Abe, however, is saying no. "What—you guys want to leave a sure thing?" he keeps asking us. "You think Microsoft's going to *shrink*, or are you nuts?"

"That's not the point, Abe."

"What *is* the point, then?"

"One-Point-Oh," I said.

"What?" replied Abe.

"Being One-Point-Oh. The first to do something cool or new."

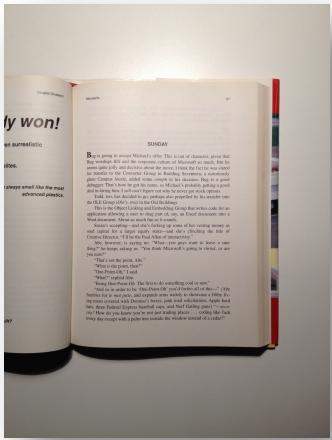
"And so in order to be 'One-Point-Oh' you'd forfeit all of this—" (Abe

fumbles for *le mot juste*, and expands arms widely to showcase a filthy liv-

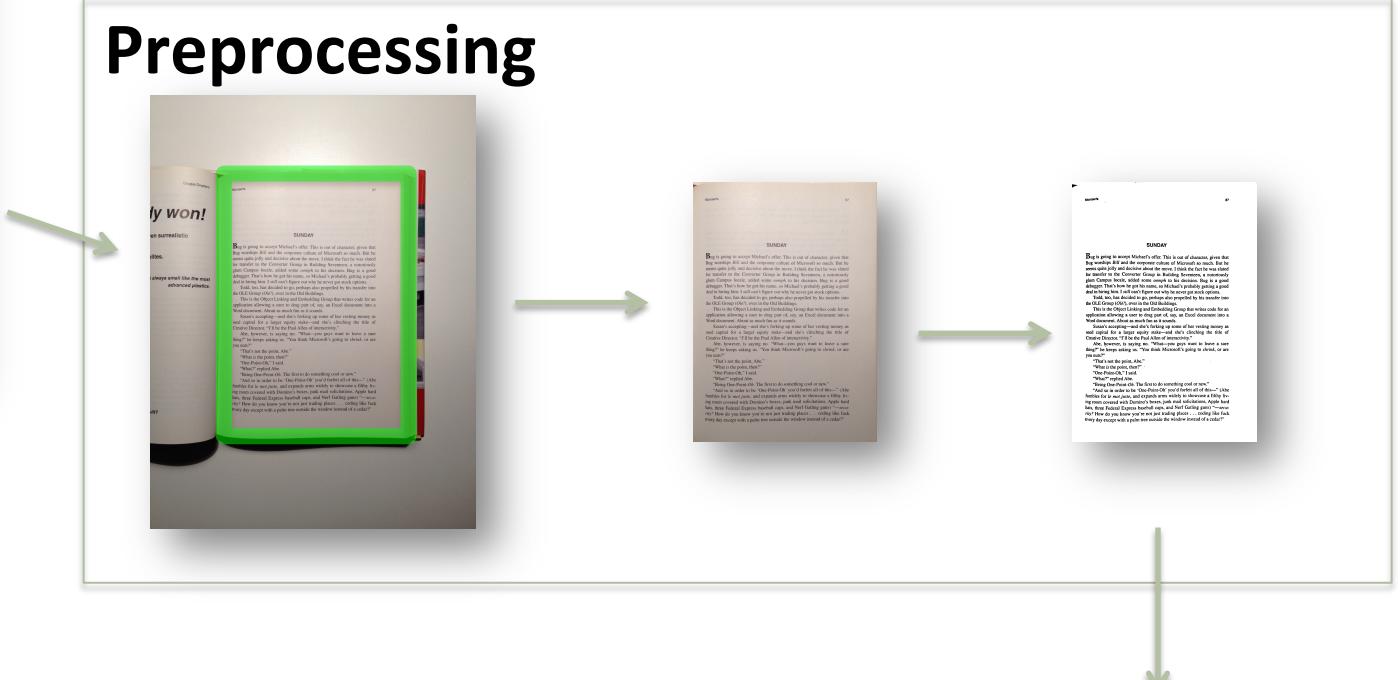
ing room covered with Domino's boxes, junk mail solicitations, Apple hard

Acquisition

OCR Pipeline

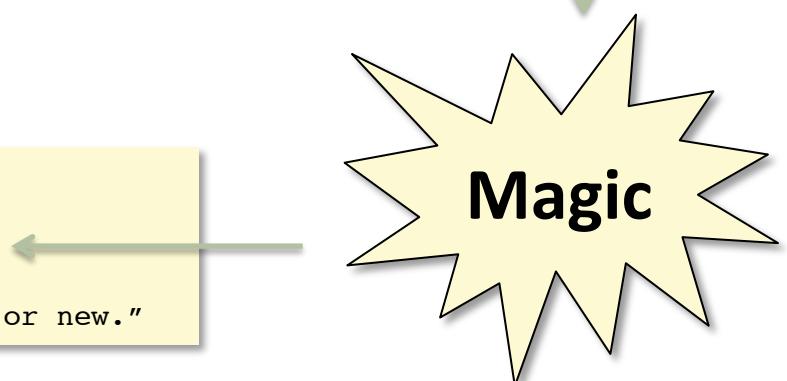


Preprocessing



Text output

"That's not the point, Abe."
"What *is* the point, then?"
"One-Point-Oh," I said.
"What?" replied Abe.
"Being One-Point-Oh. The first to do something cool or new."



OpenCV – Computer Vision Library

- 2500 CV Algorithms
- BSD license
- All main platforms
- CUDA, OpenCL acceleration
- Realtime, multi-core
- <http://opencv.org>

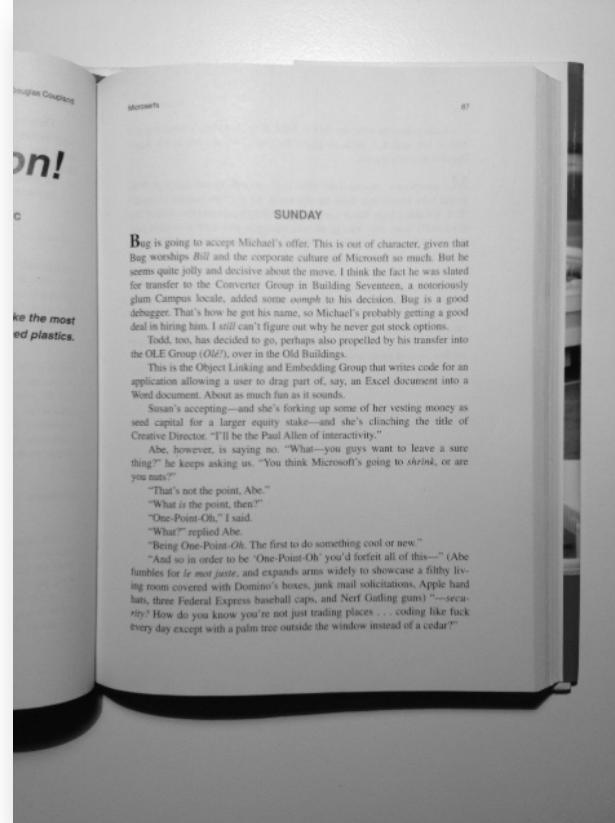


	OpenCV for Windows
	OpenCV for Linux/Mac
	OpenCV for Android
	OpenCV for iOS

Where is the Page?

Trick 1: Reduce Size and Colours

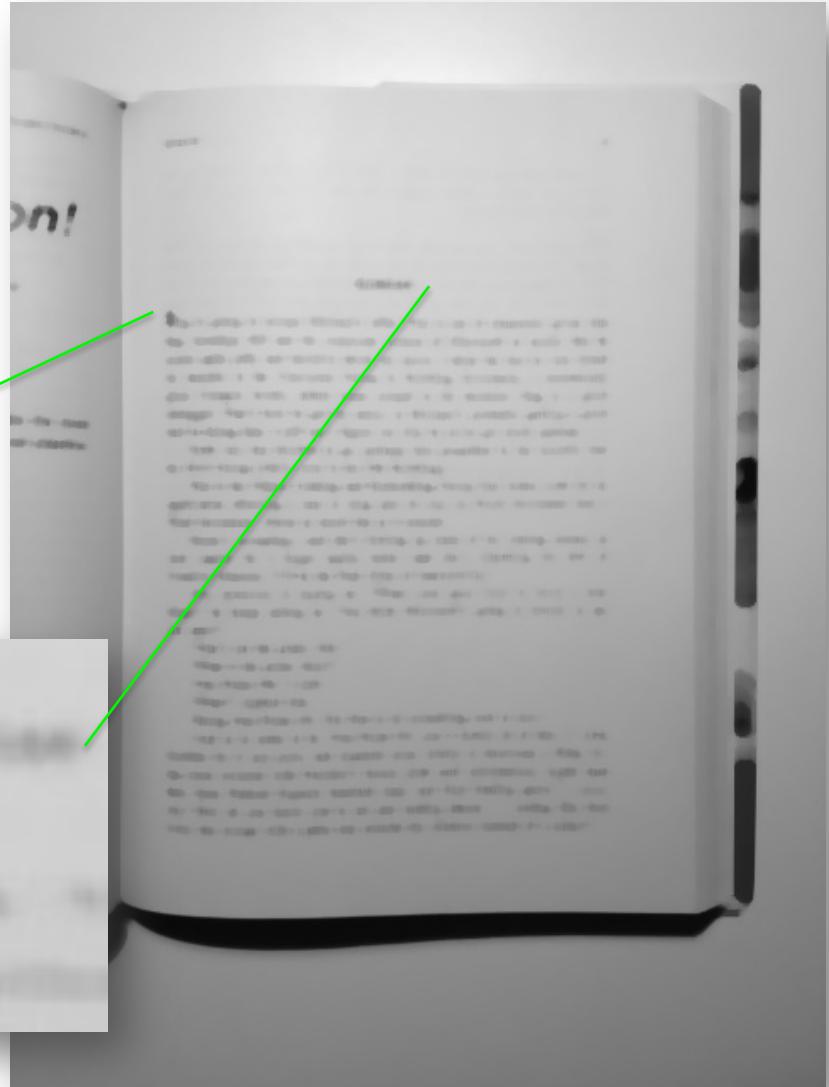
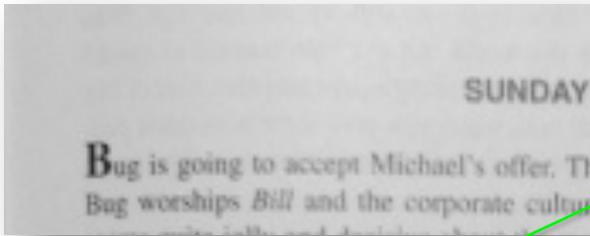
```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    show('1: Input image', gray)
```



Where is the Page?

Trick 2: Remove Noise and Details

```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    show('1: Input image', gray)
    # Remove high-frequency noise (e.g. letters)
    median = cv2.medianBlur(gray, 7)
    show('2: Median filtered', median)
```

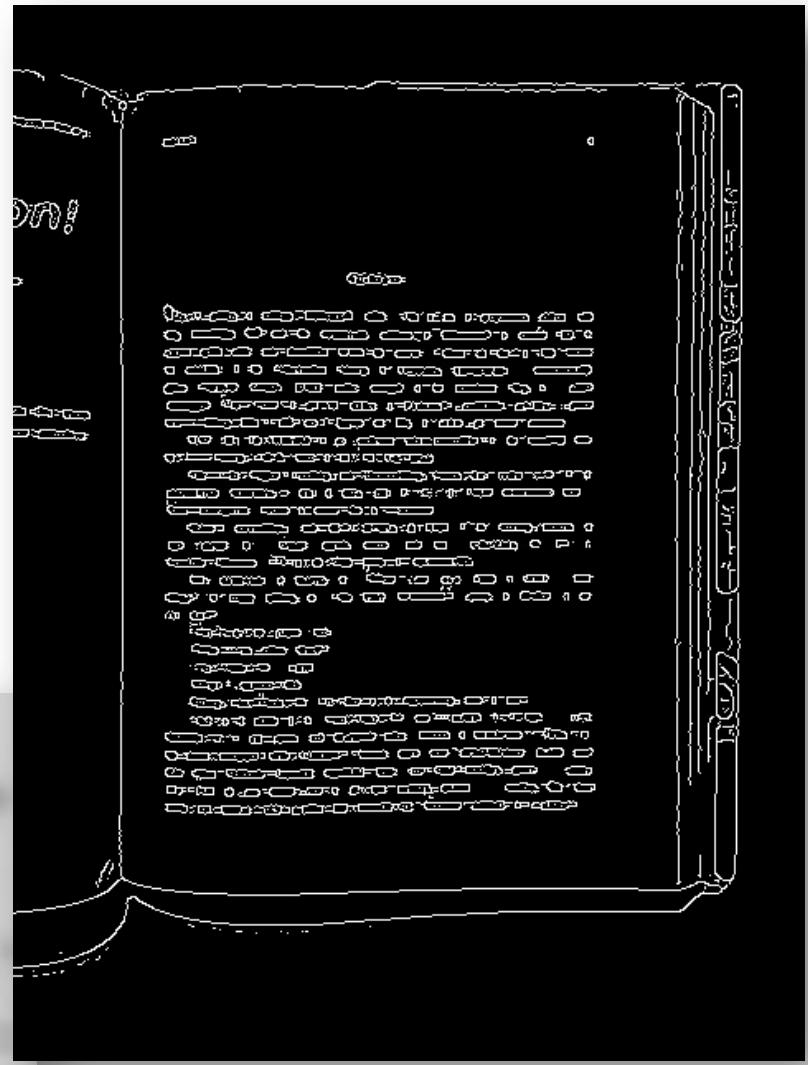


Trick 3: Find Edges

```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    show('1: Input image', gray)
    # Remove high-frequency noise (e.g. letters)
    median = cv2.medianBlur(gray, 7)
    show('2: Median filtered', median)

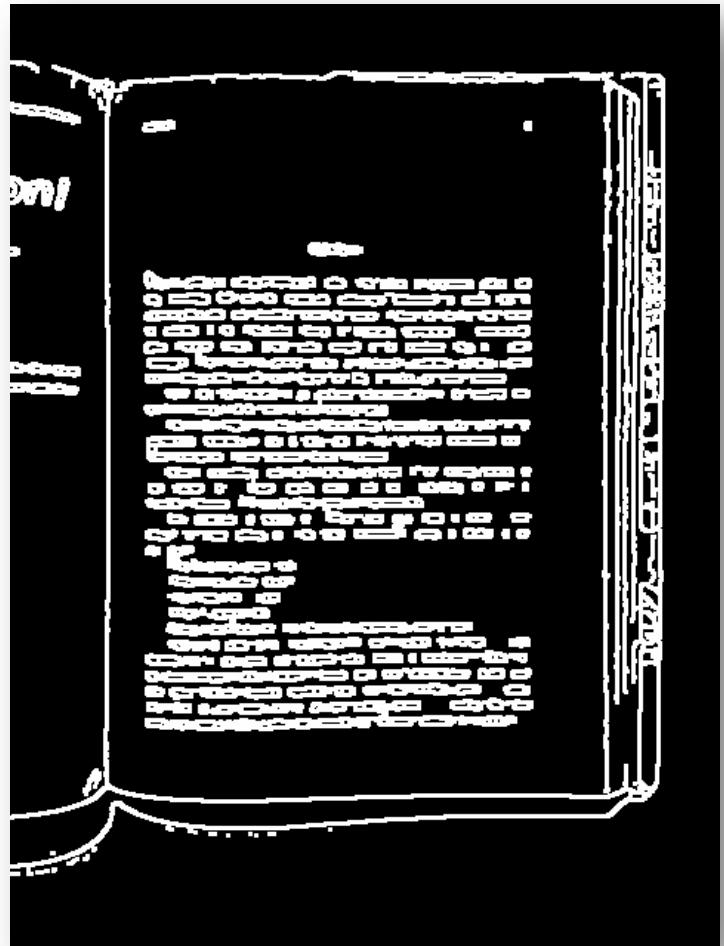
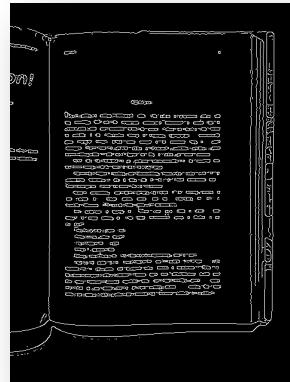
    edges = cv2.Canny(median, 10, 20,
                      apertureSize=3)

    show('3: Edges', edges)
```



Trick 4: Dilate to Thicken

```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    show('1: Input image', gray)
    # Remove high-frequency noise (e.g. letters)
    median = cv2.medianBlur(gray, 7)
    show('2: Median filtered', median)
    edges = cv2.Canny(median, 10, 20, apertureSize=3)
    show('3: Edges', edges)
    dilated = cv2.dilate(edges,
                         kernel=None,
                         iterations=1)
    show('4: Dilated Edges', dilated)
```



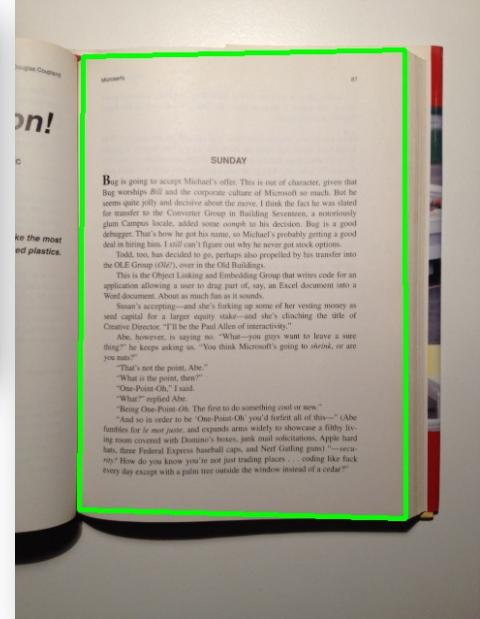
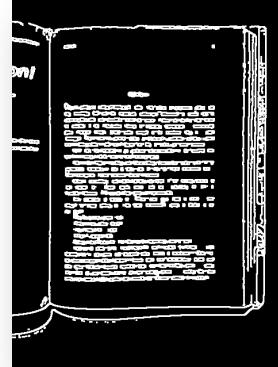
Example code: [src/books.py](#)

Trick 5: Find Contours

```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    show('1: Input image', gray)
    # Remove high-frequency noise (e.g. letters)
    median = cv2.medianBlur(gray, 7)
    show('2: Median filtered', median)
    edges = cv2.Canny(median, 10, 20, apertureSize=3)
    show('3: Edges', edges)
    dilated = cv2.dilate(edges, kernel=None, iterations=1)
    show('4: Dilated Edges', dilated)

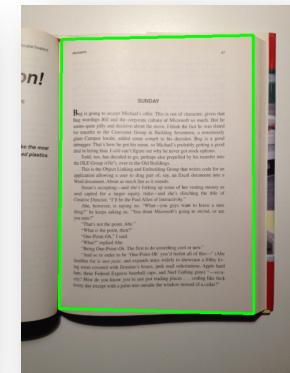
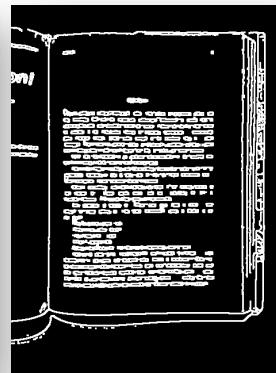
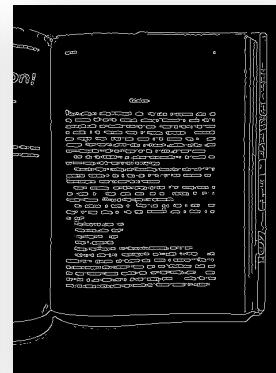
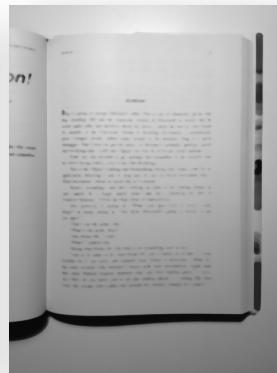
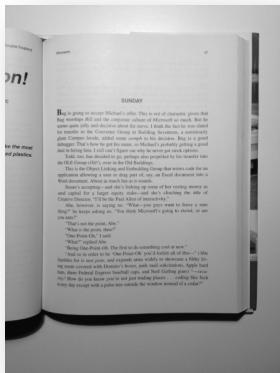
    min_page_area = img.shape[0] * img.shape[1] / 4
    quads = find_quadrilaterals(dilated, min_area=min_page_area, cos_epsilon=.8)
    quads_found = draw_contours(img, quads)
    show('5: Page found', quads_found)

def find_quadrilaterals(img, min_area, cos_epsilon):
    quads = []
    contours, hierarchy = cv2.findContours(img.copy(), cv2.RETR_LIST, cv2.CHAIN_APPROX_SIMPLE)
    for cnt in contours:
        cnt_len = cv2.arcLength(cnt, True)
        poly = cv2.approxPolyDP(cnt, 0.02*cnt_len, True)
        if len(poly) == 4 and cv2.contourArea(poly) > min_area and cv2.isContourConvex(poly):
```



Where is the Page?

```
def find_page_with_canny(img):
    '''Find the page in the image.'''
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    median = cv2.medianBlur(gray, 7)
    edges = cv2.Canny(median, 10, 20, apertureSize=3)
    dilated = cv2.dilate(edges, kernel=None, iterations=1)
    min_page_area = img.shape[0] * img.shape[1] / 4
    quads = find_quadrilaterals(dilated, min_area=min_page_area, cos_epsilon=.8)
    quads_found = draw_contours(img, quads)
```



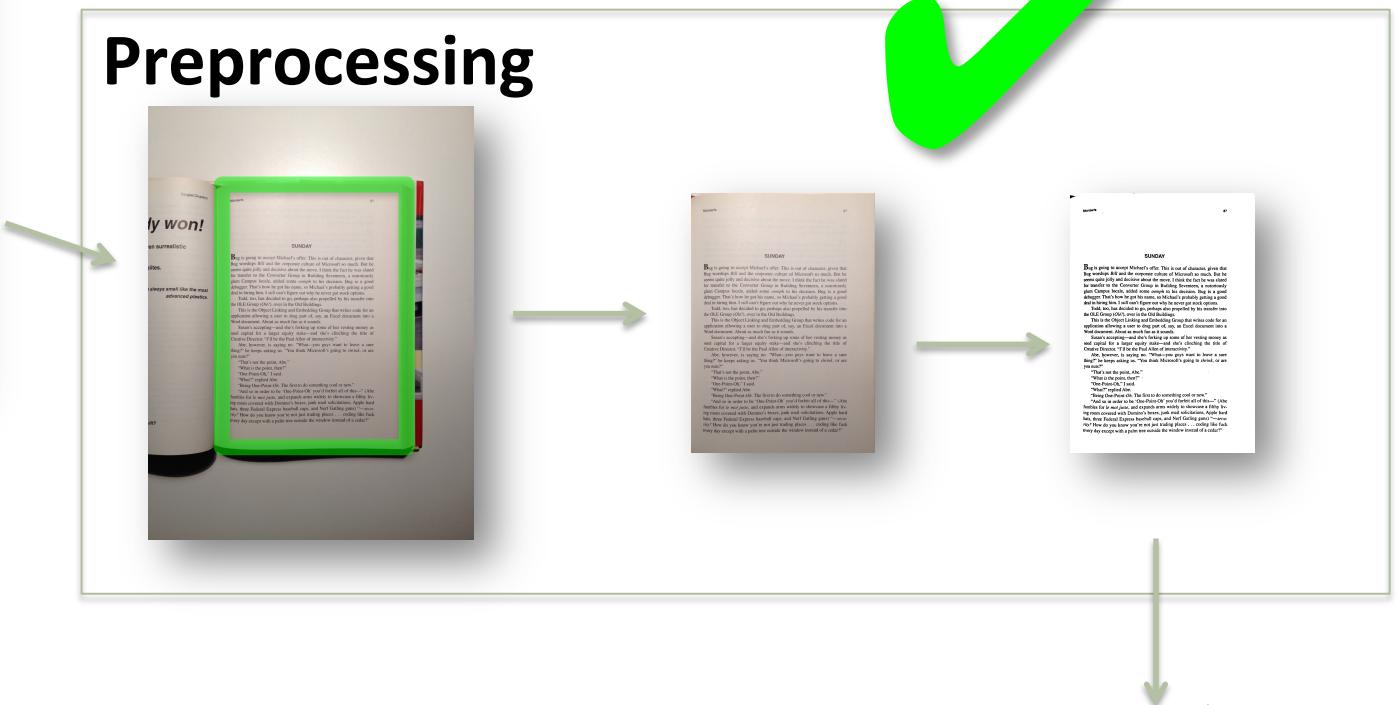
Example code: [src/books.py](#)

Acquisition

OCR Pipeline

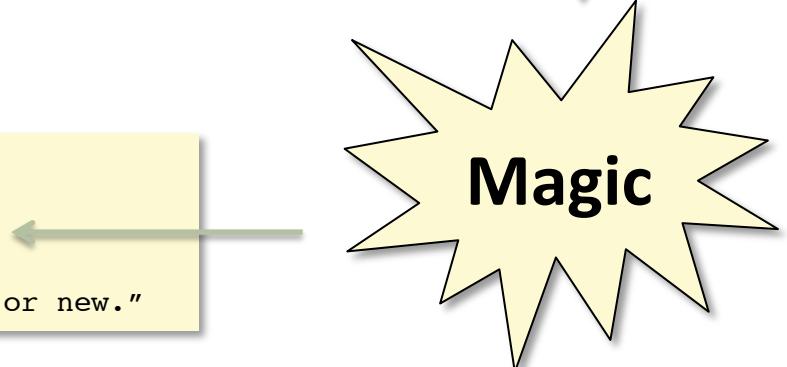


Preprocessing

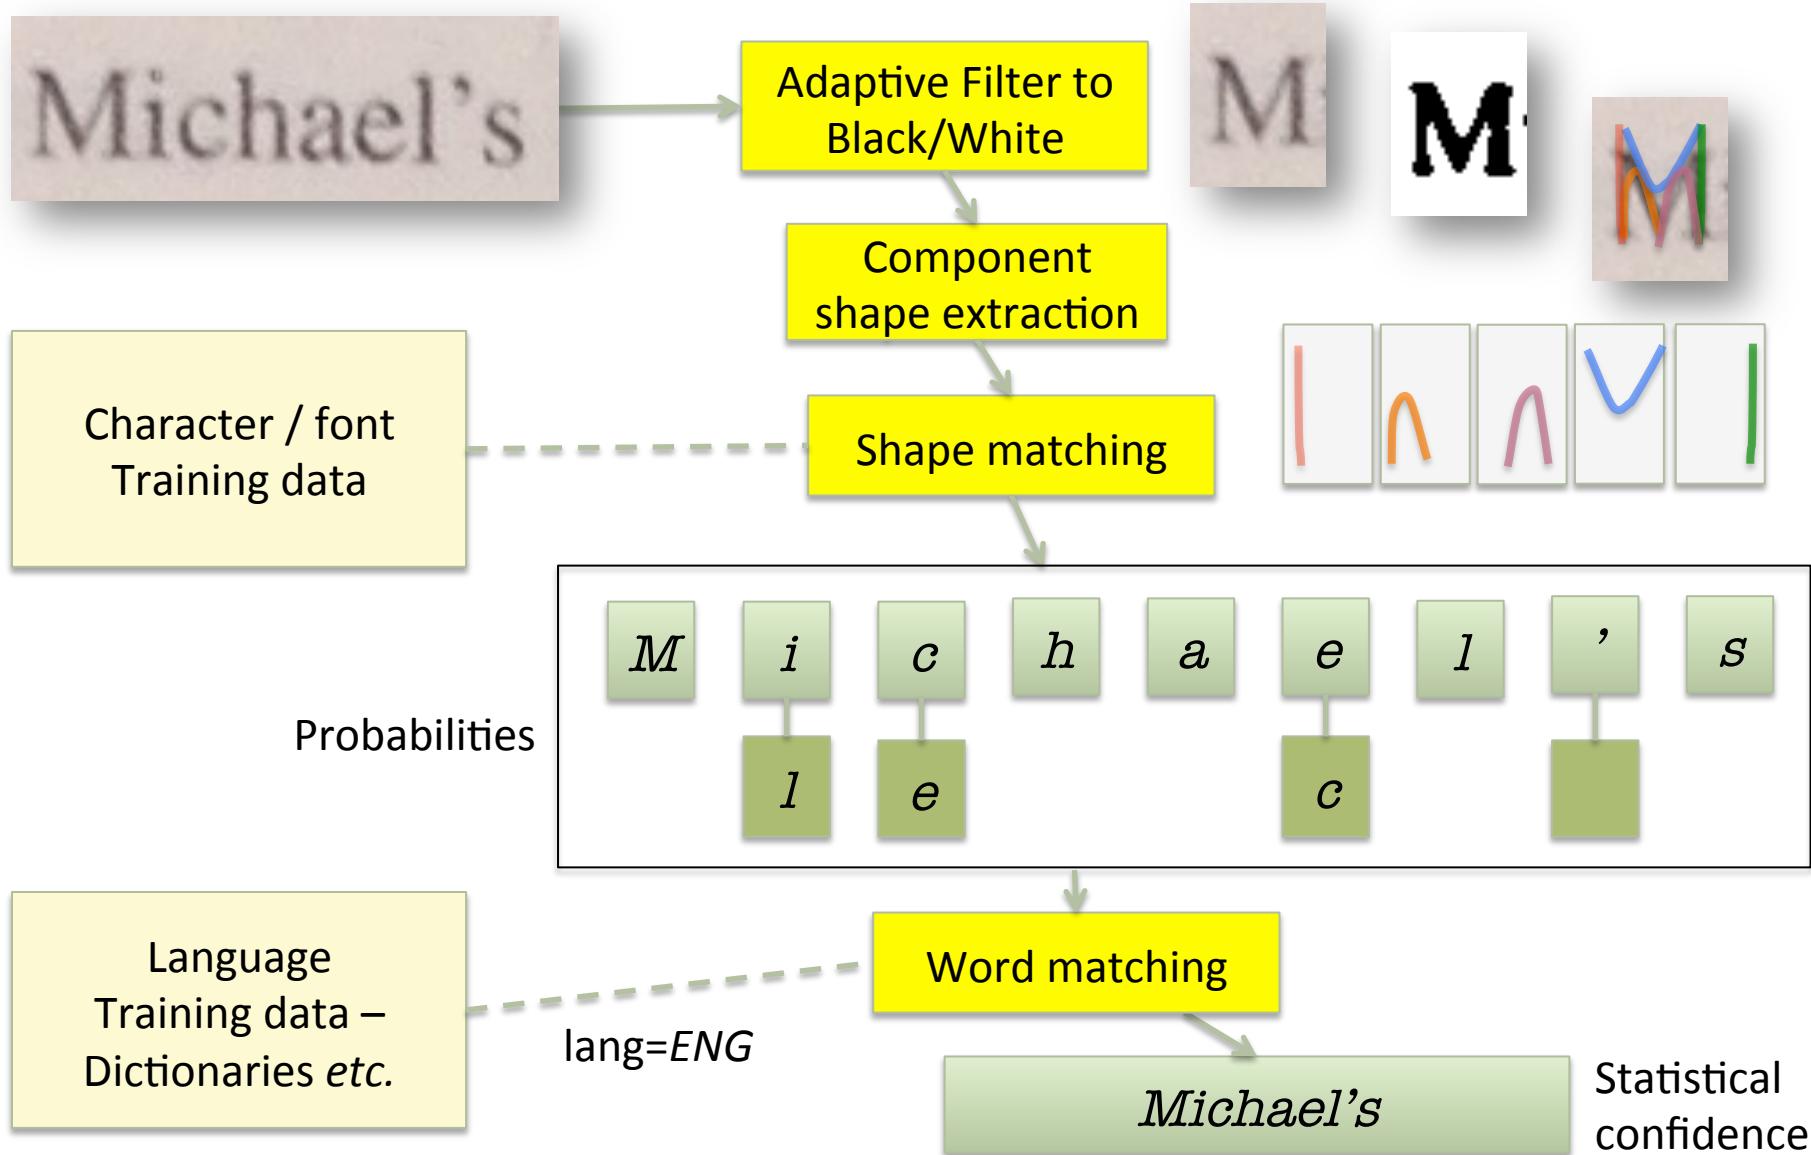


Text output

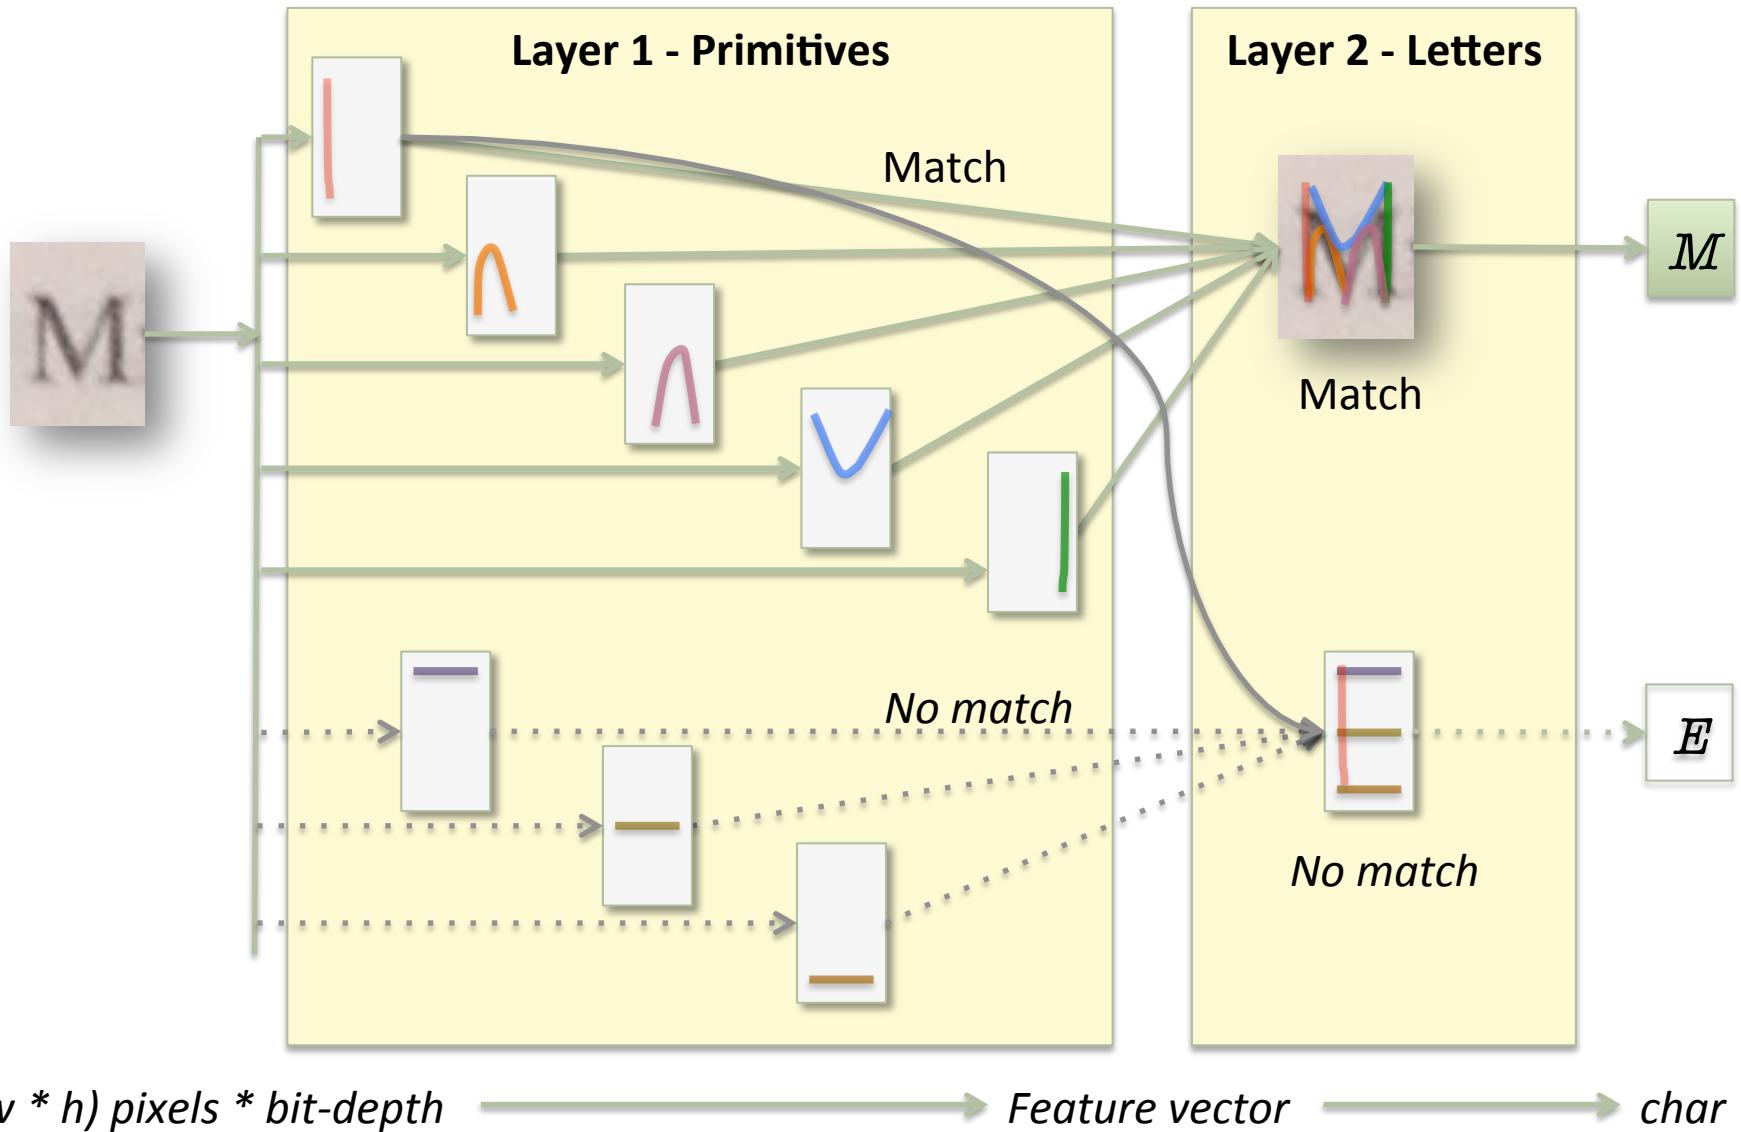
"That's not the point, Abe."
"What *is* the point, then?"
"One-Point-Oh," I said.
"What?" replied Abe.
"Being One-Point-Oh. The first to do something cool or new."



Magic Explained



OCR Shape Matching



Tesseract OCR Library

Free, open-source software

Developed by HP Labs 1985-1995

Now sponsored by Google



- C++
- Windows, Mac, Linux, Android, iOS
- Wrappers for .NET, Python, ...
- Apache License 2.0
- <http://code.google.com/p/tesseract-ocr/>

1987 - HP introduces the desktop scanner, ScanJet
http://www.hpmuseum.net/display_item.php?hw=295

OCR - Magic Happens Here

```
def ocr_text(img):  
    '''Perform OCR on the image.'''  
    tr = Tesseract(lang='eng')  
    pil_image = pil.Image.fromarray(img)  
    tr.set_image(pil_image)  
    utf8_text = tr.get_text()  
    return utf8_text
```

SUNDAY

Bug is going to accept Michael's offer. This is out of character, given that Bug worships Bill and the corporate culture of Microsoft so much. But he seems quite jolly and decisive about the move. I think the fact he was slated for transfer to the Converter Group in Building Seventeen, a notoriously

SUNDAY

Bug is going to accept Michael's offer. This is out of character, given that Bug worships Bill and the corporate culture of Microsoft so much. But he seems quite jolly and decisive about the move. I think the fact he was slated for transfer to the Converter Group in Building Seventeen, a notoriously

Application: Feature Extraction

CARS

Pipeline



Find Plate



Read Plate

VX 59 134

Look up plate

Danish Police (2013)

400,000 “wanted” plates in Denmark

Automatic scanning and lookup from police cars

Vision pilot project: 20 million DKK (22 M NOK)

(Norway already has this)

Find the Plate

- Find rectangles in multiple representations
 - Find edges at various grayscale levels (threshold to B/W)
 - Detect edges, extract rectangles as before
 - Plates have (approximately) 5:1 proportions



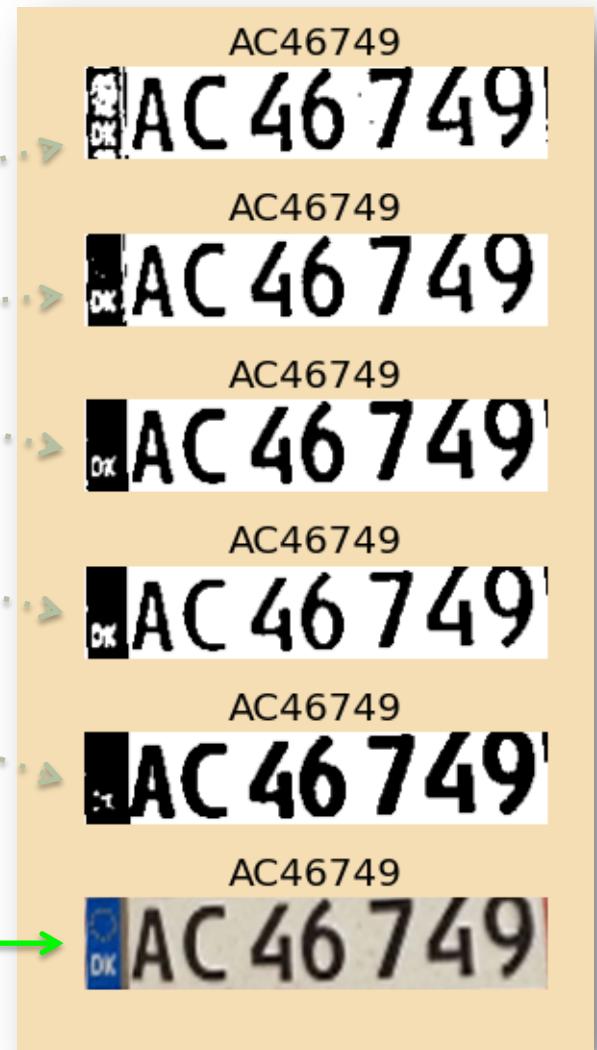
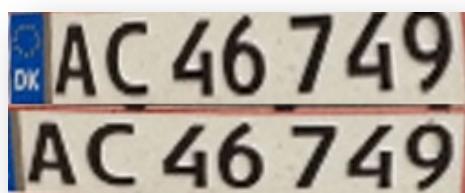
Two candidate plate regions found



Example code: src/cars.py

Read the Plate

- Transform perspective to “head on”
- Filter plate to different representations
 - Adaptive threshold
 - Manually threshold at middle gray
 - Black/white with optimal threshold (“OTSU”)
 - Erode/dilate letters
- OCR all
- Match OCR texts to “2 letters, 5 digits”
- Select most frequently matched



Example code: `src/cars.py`

Plate Filtering

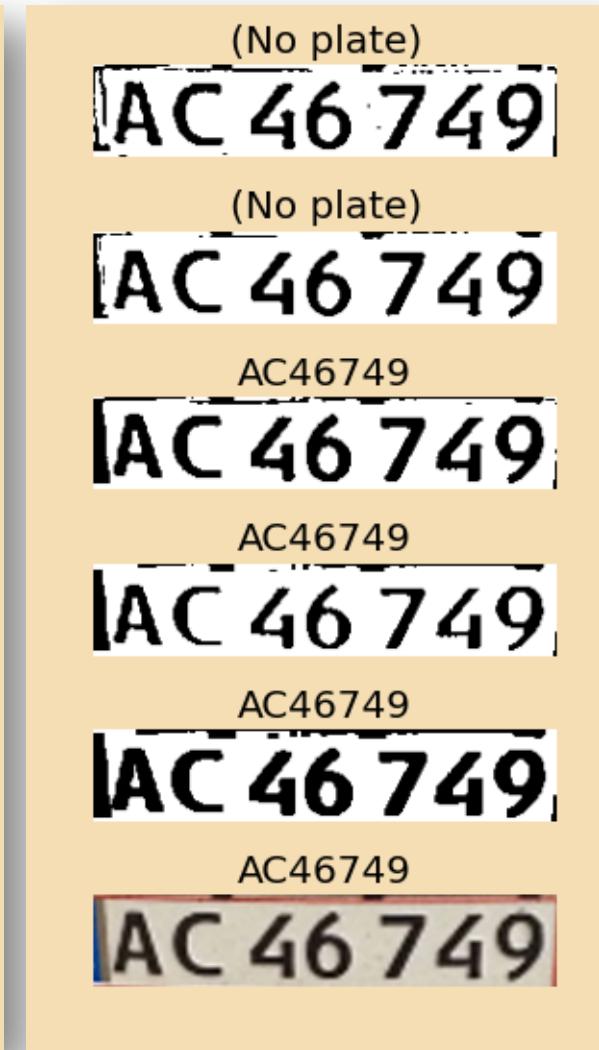
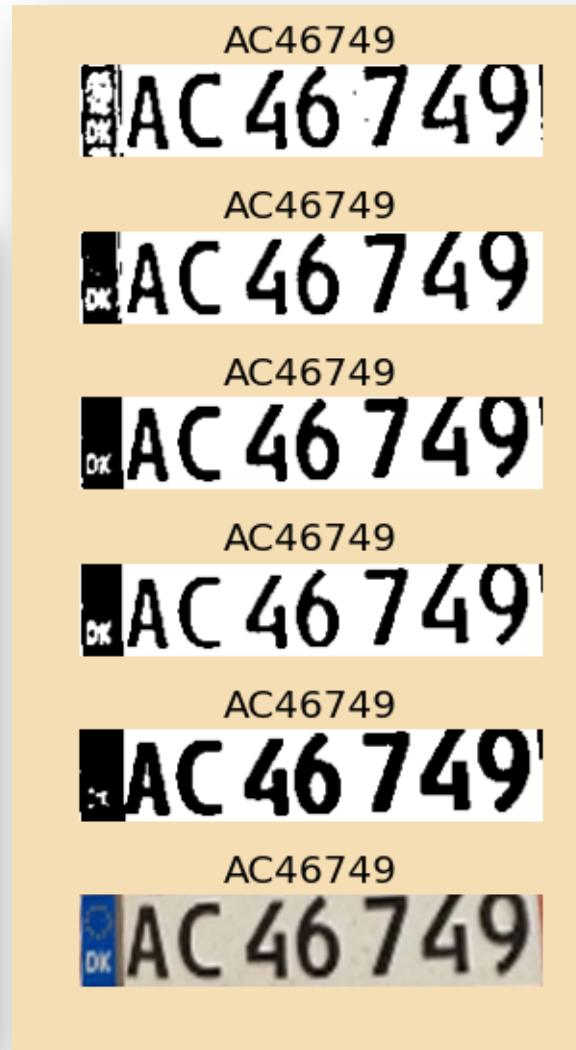
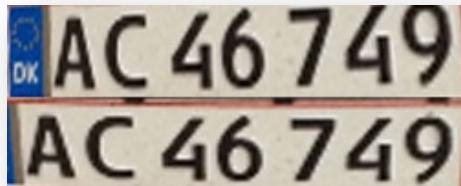
```
def match_plates(title, candidate_plate_images):
    '''Returns a list of the possible matches for the plate.'''
    matches = []
    n = 0
    for cp in candidate_plate_images:
        # Filter to different representations
        n = n+1
        gray = cv2.cvtColor(cp, cv2.COLOR_BGR2GRAY)
        adaptive = cv2.adaptiveThreshold(gray, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
                                         cv2.THRESH_BINARY, 11, 15)
        ret_val, th = cv2.threshold(gray, 100, 255, cv2.THRESH_BINARY)
        t_otsu, th_otsu = cv2.threshold(gray, 100, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)
        ret_val, th_otsu = cv2.threshold(gray, t_otsu, 255, cv2.THRESH_BINARY)
        # thin the black parts (letters and noise)
        dilated = cv2.dilate(th_otsu, kernel=(5,5), iterations=2)
        # thicken the letters
        eroded = cv2.erode(dilated, None)
        image_variants = [cp, adaptive, th, th_otsu, dilated, eroded]
        image_matches = [m for m in map(ocr_plate, image_variants)]
        if PLOT_MATCHES:
            plot_image_variants_and_matches("%s matches %d" % (title, n),
                                             image_variants, image_matches)
        matches += [m for m in image_matches if m]
    return matches
```



OCR Code

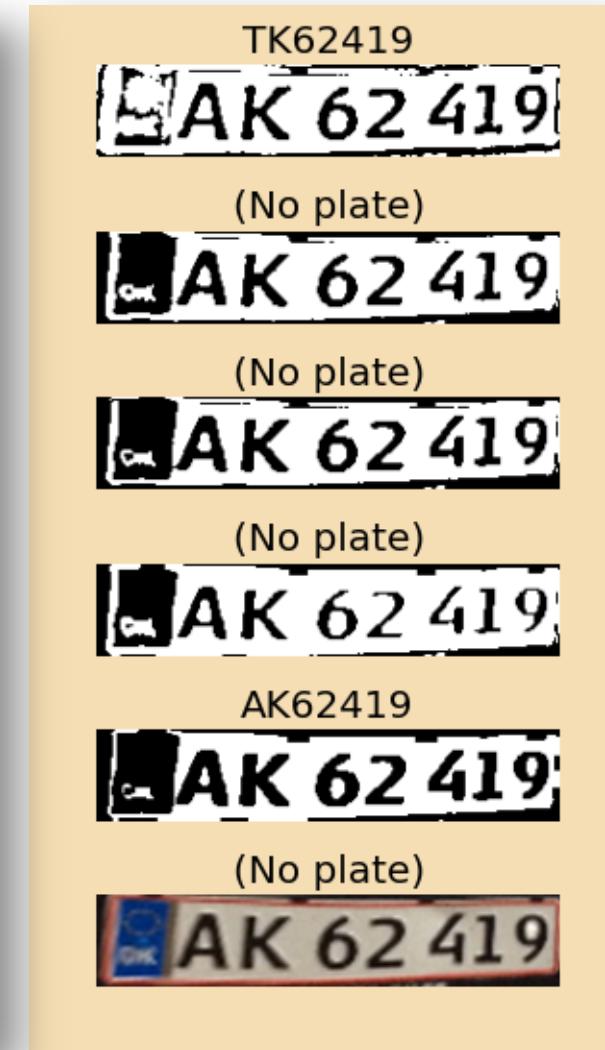
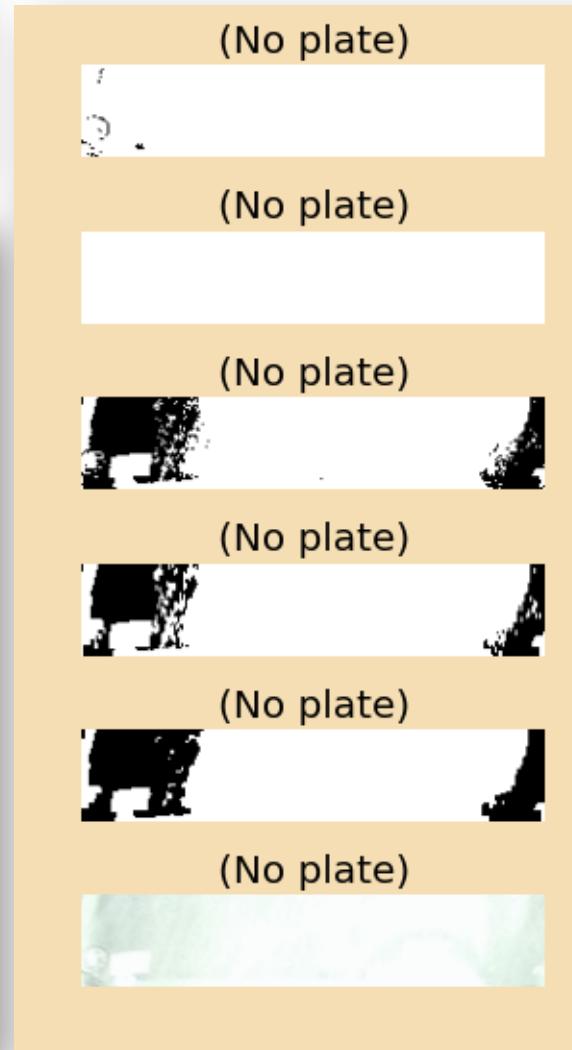
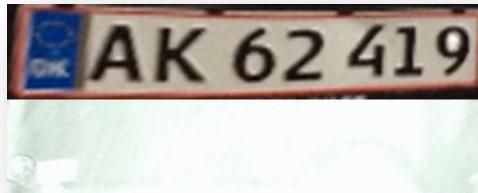
```
# -----
# Text recognition
# -----  
  
def ocr_text(img):  
    tr = Tesseract(lang='eng')  
    tr.clear()  
    pil_image = pil.Image.fromarray(img)  
    # Turn off OCR word dictionaries  
    tr.set_variable('load_system_dawg', "F")  
    tr.set_variable('load_freq_dawg', "F")  
    tr.set_variable('-psm', "7") # treat image as single line  
    tr.set_variable('tessedit_char_whitelist', "ABCDEFGHIJKLMNOPQRSTUVWXYZ0123456789")  
    tr.set_image(pil_image)  
    utf8_text = tr.get_text()  
    return unicode(utf8_text)  
  
def ocr_plate(img):  
    text = ocr_text(img)  
    match = re.search(r"[A-Z][A-Z] *\d{2} *\d{3}", text)  
    result = None  
    if match:  
        result = match.group()  
        # Canonicalize to no spacing  
        result = result.replace(' ', '')  
    return result
```

Number Plate Recognition



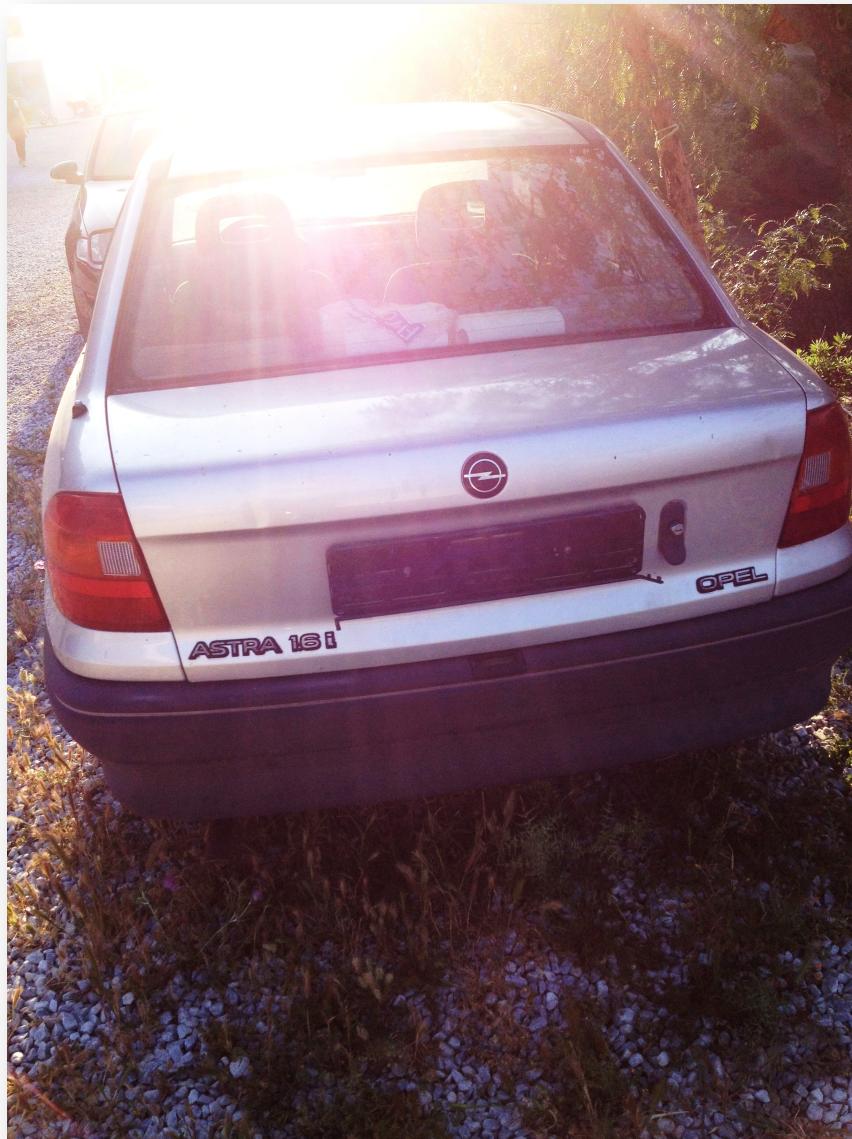
Example code: [src/cars.py](#)

Number Plate Recognition



Example code: [src/cars.py](#)

#FAIL



Feature Recognition

PIZZA

How would you do it?



Pizza
time!



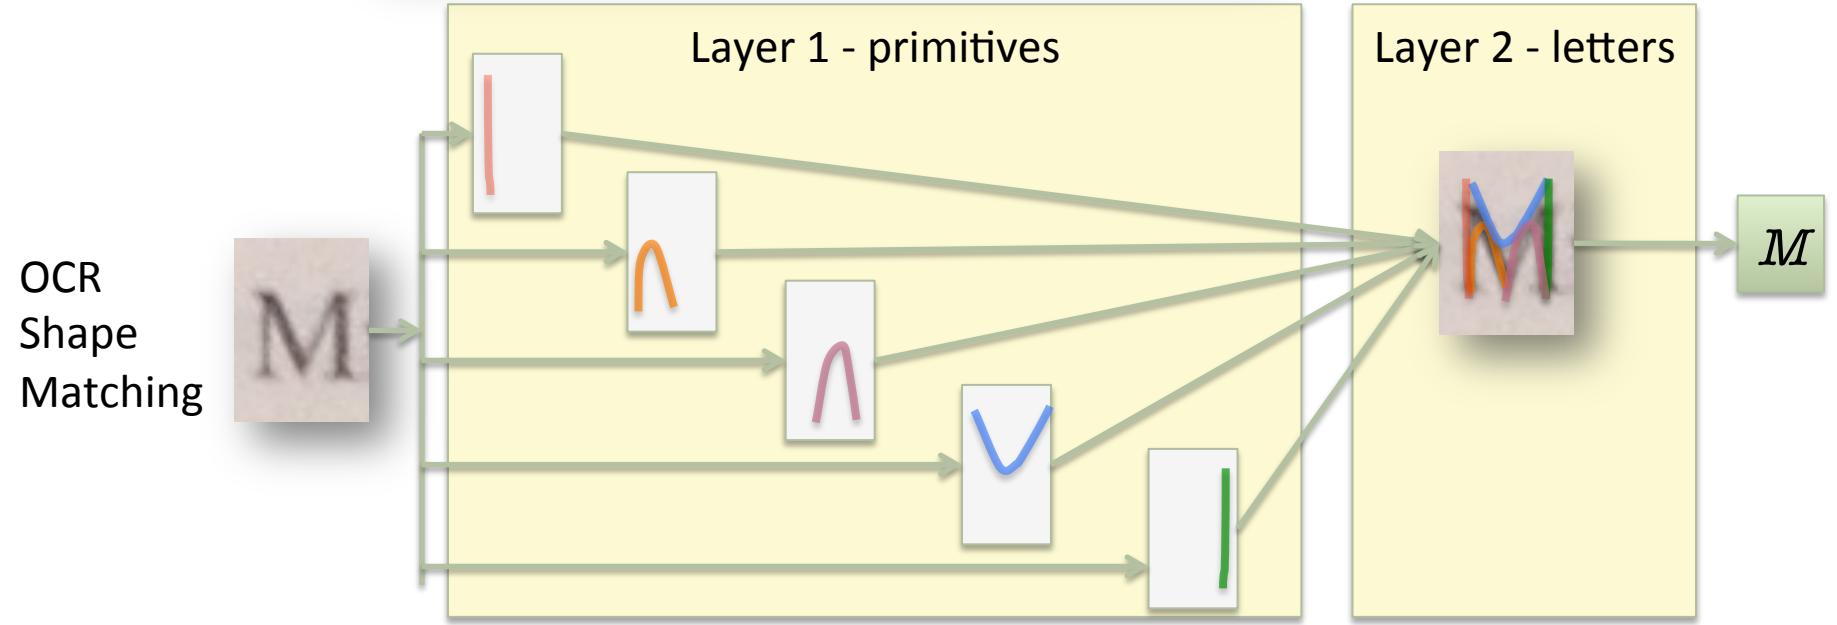
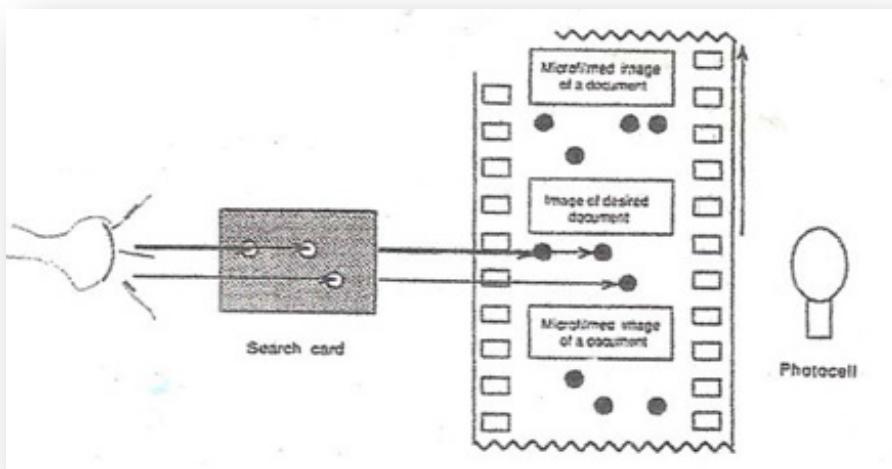
20 Questions...

- Quadrilateral found?
 - Is it dominantly white?
 - Percentage of surface
 - Black
 - Orange
 - Grey
 - Skin tones on edge of quad?
 - *Is there a face shape?*
 - *Is face above and near quad?*
 - ...
- Probably a pizza box



1920s Revisited

1920s
Template matching



Known Object Detection



Scale



Position



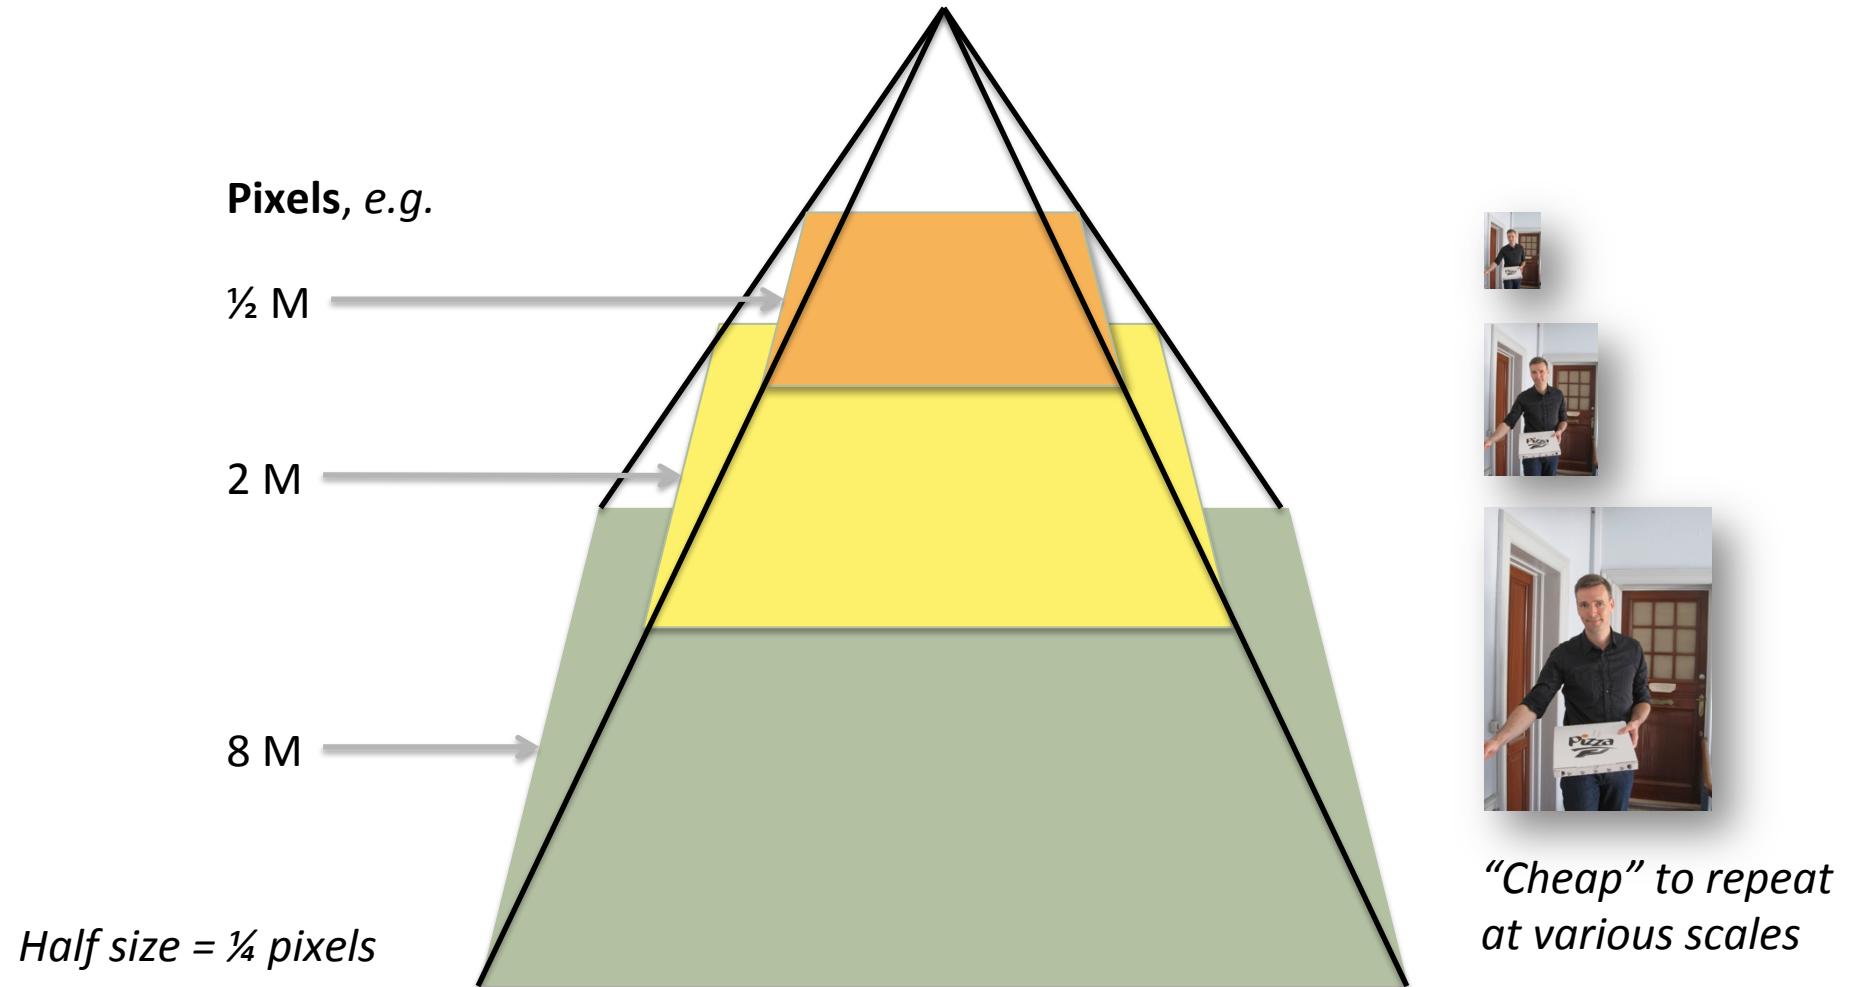
Brightness /
illumination



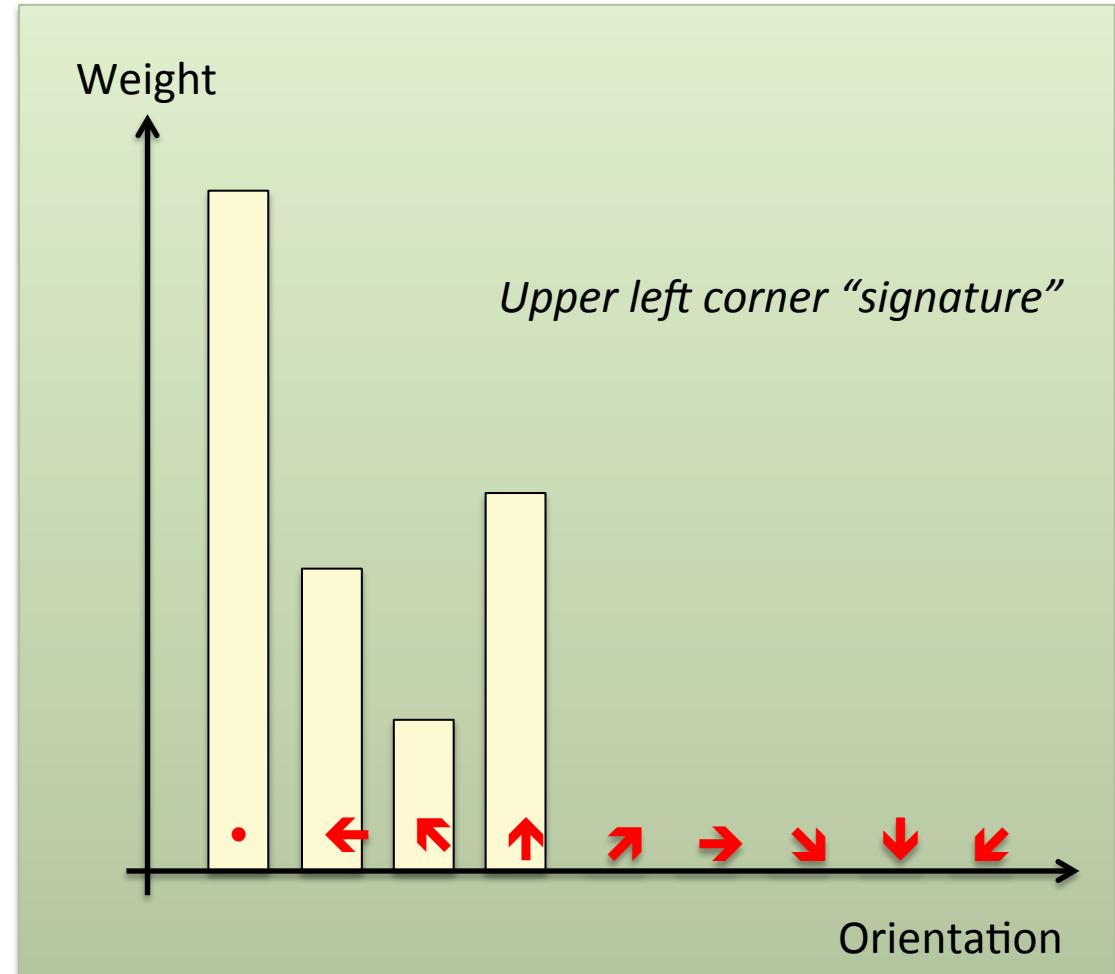
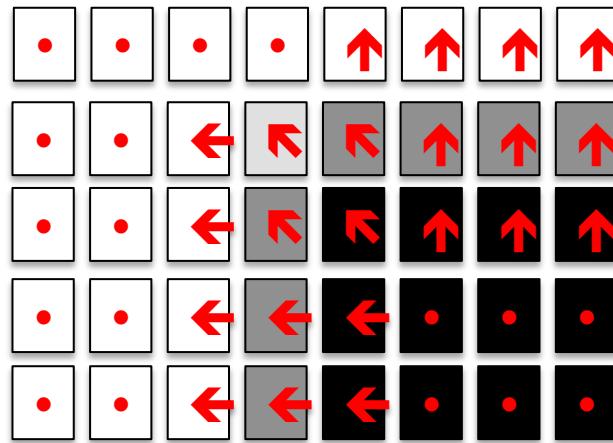
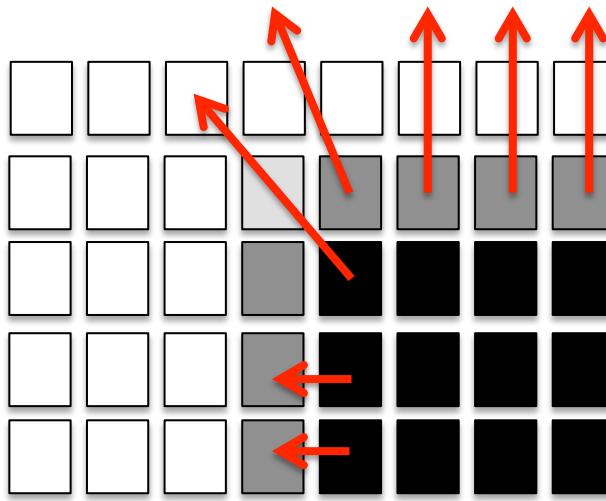
Orientation



Scale: Pyramid or Scale-Space



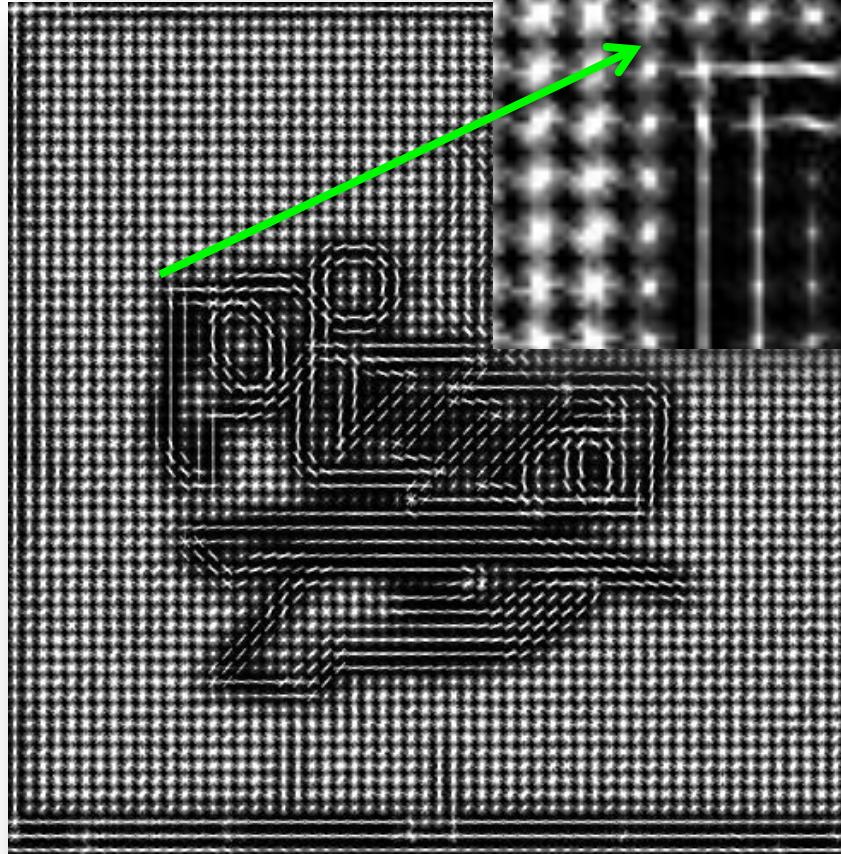
HOG – Histogram of Oriented Gradients



HOG Example



Input



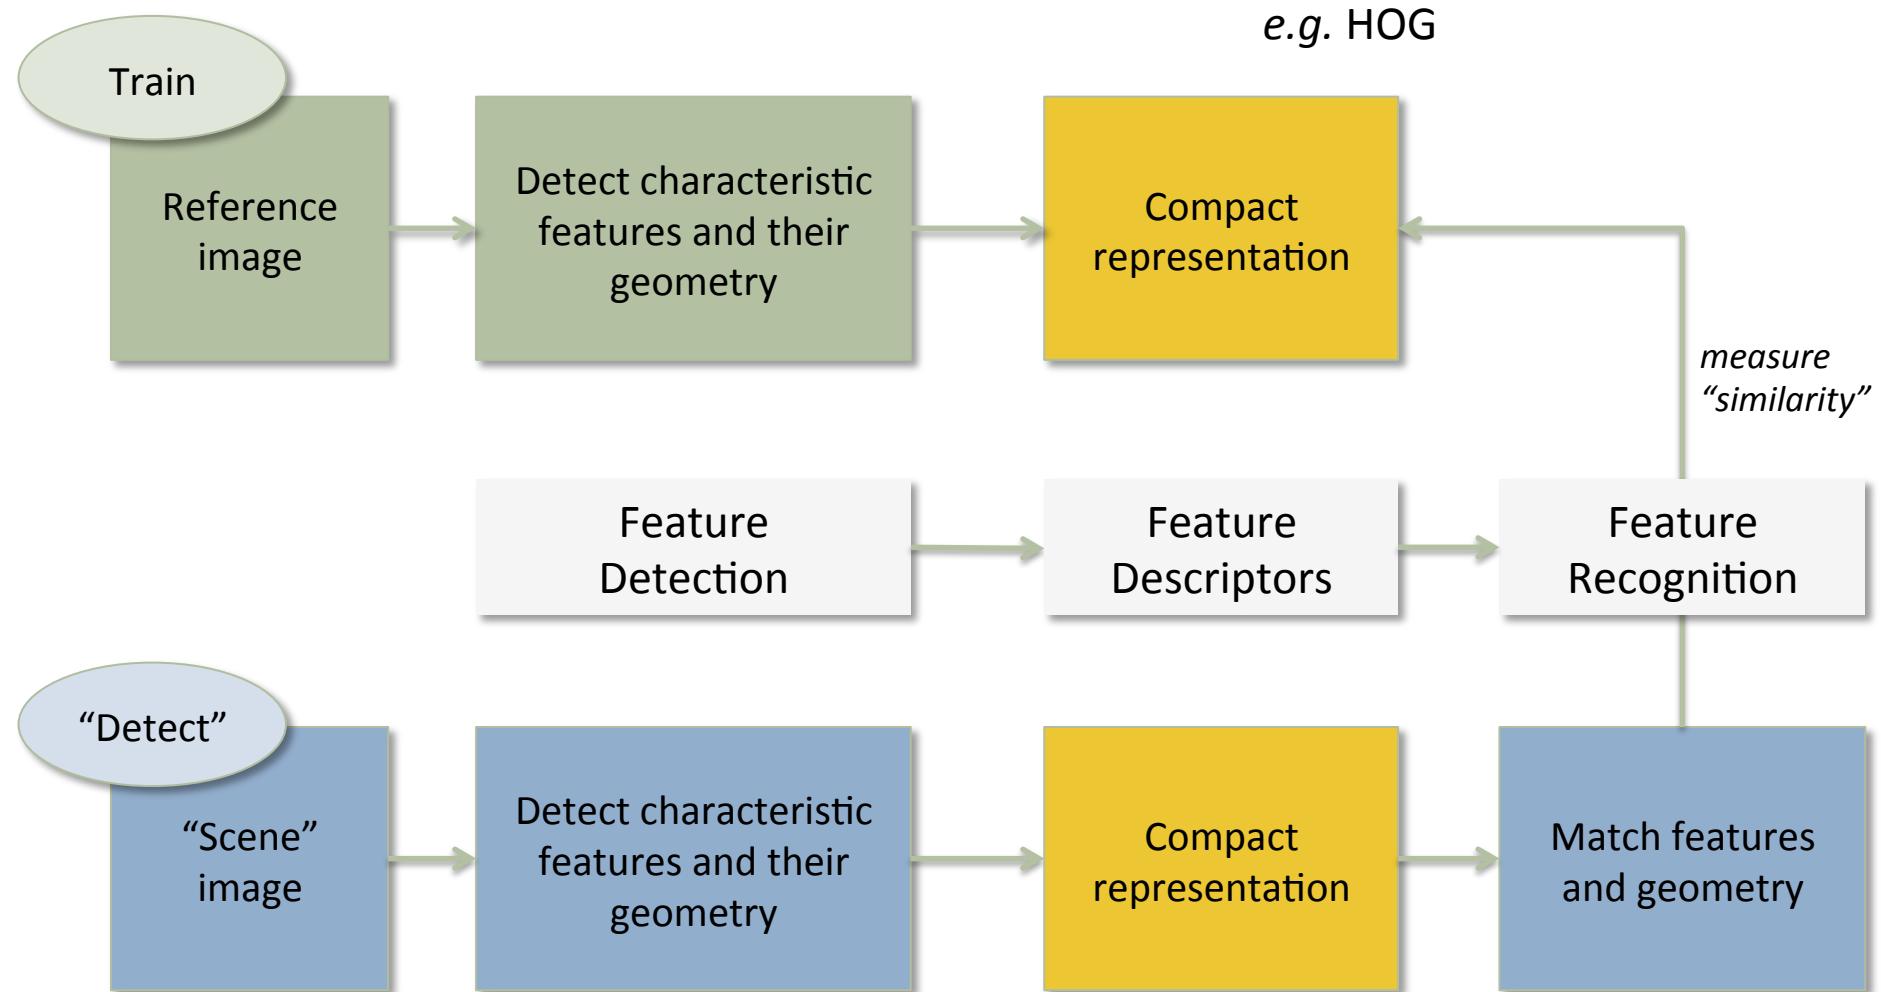
HOG



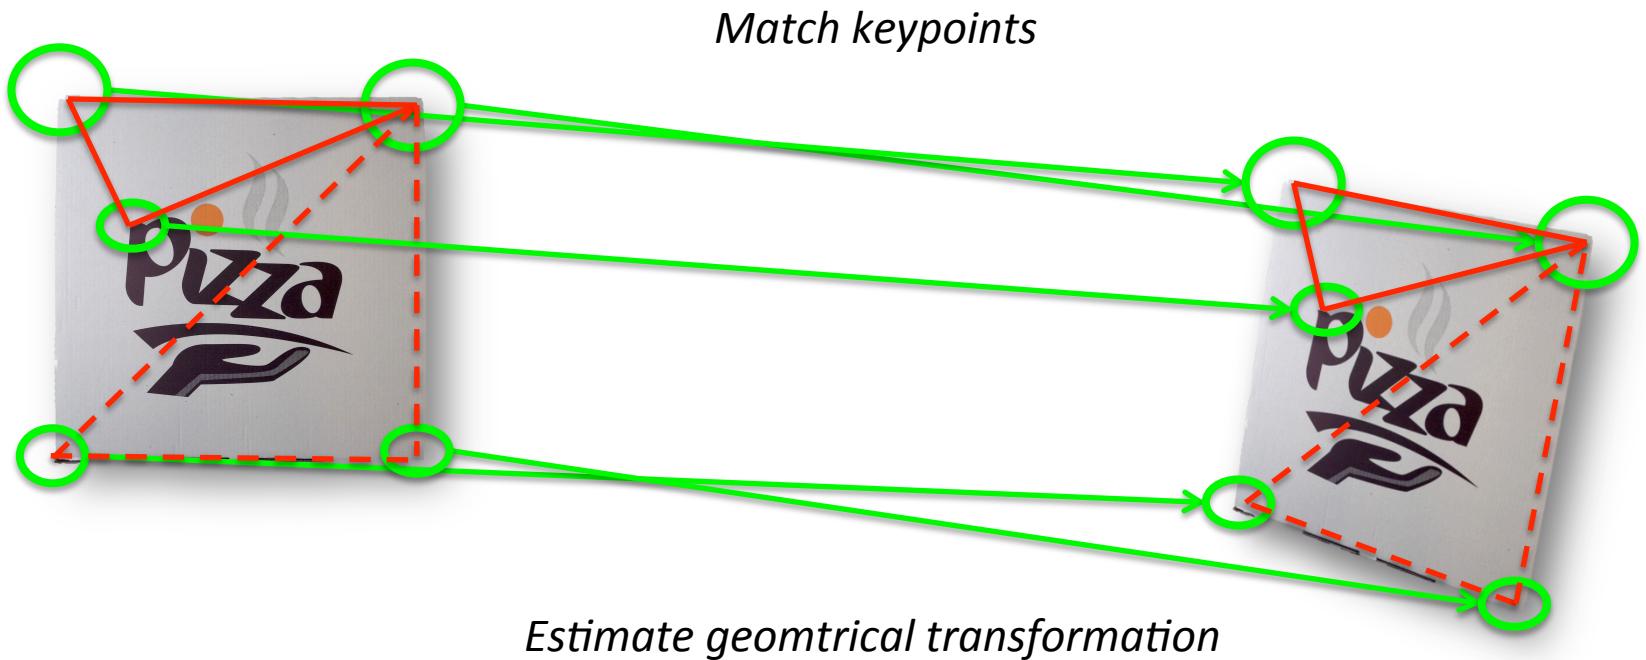
HOG reverse

Images generated with HOG visualizer at <http://web.mit.edu/vondrick/ihog/index.html>

Feature Recognition



Feature Detection and Recognition



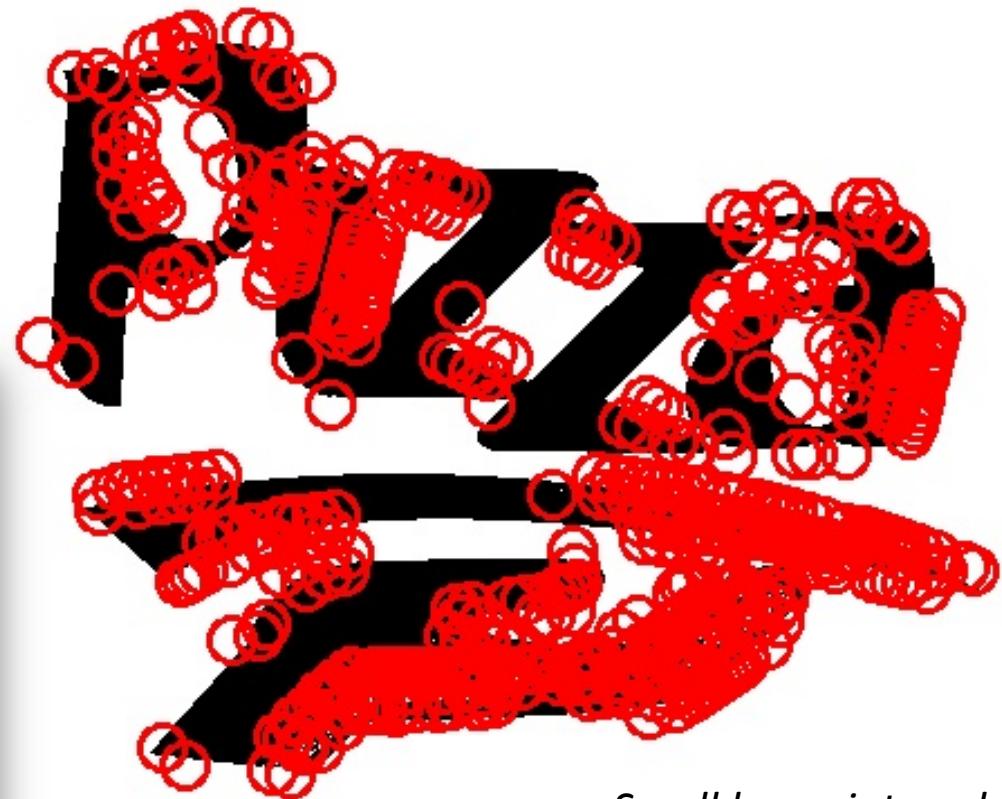
Features Detection - Keypoints



Object



All keypoints



Small keypoints only

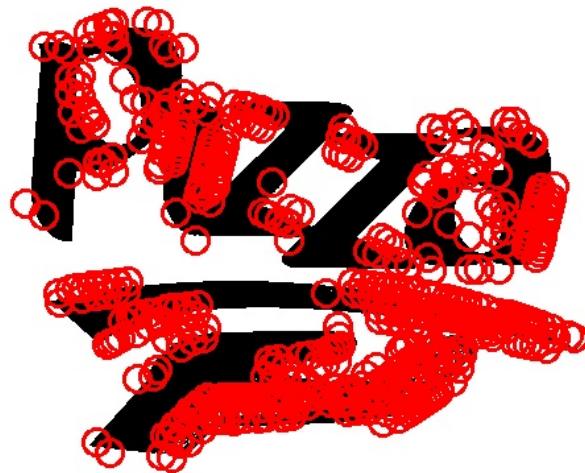
Example code: src/pizza.py

Features Detection - Keypoints



Example code: [src/pizza.py](#)

Feature Recognition



Pizza Code

```
detector, norm = cv2.SURF(), cv2.NORM_L2
bf_matcher = cv2.BFMatcher(norm)

# Detect features in the reference object
object_keypoints, object_descriptors = detector.detectAndCompute(object_image, None)

.....
.....
.....

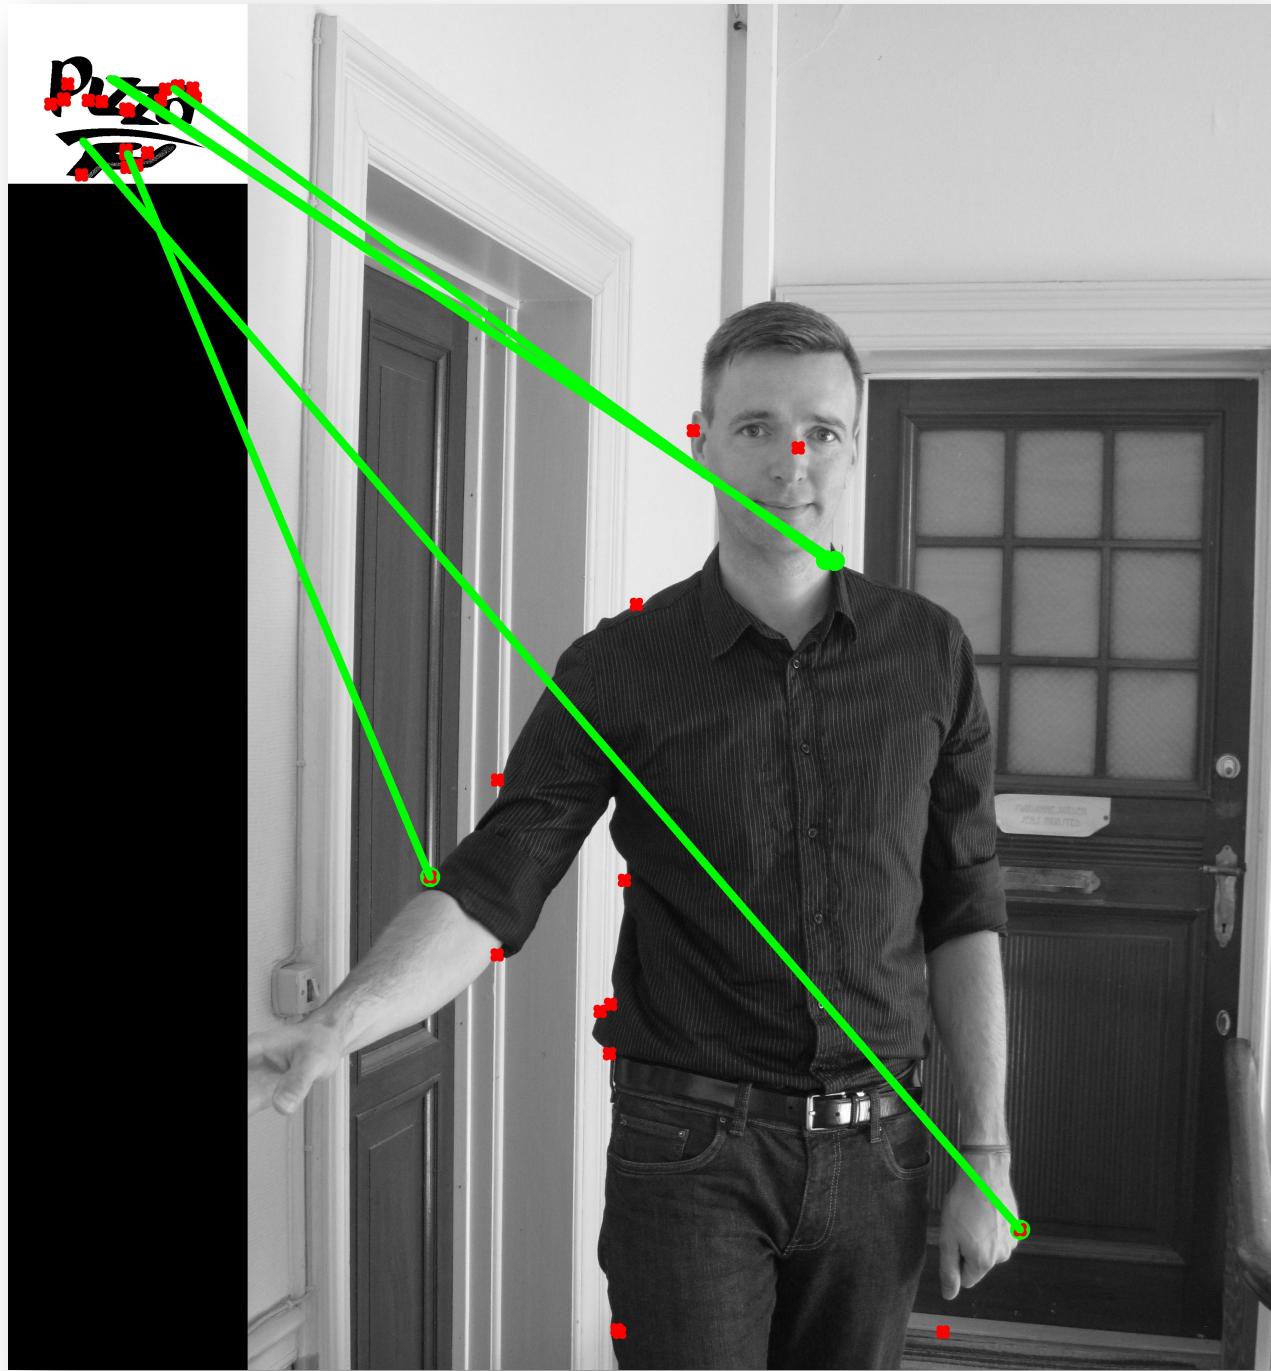


# Detect features in the scene and match
scene_keypoints, scene_descriptors = detector.detectAndCompute(scene_image, None)

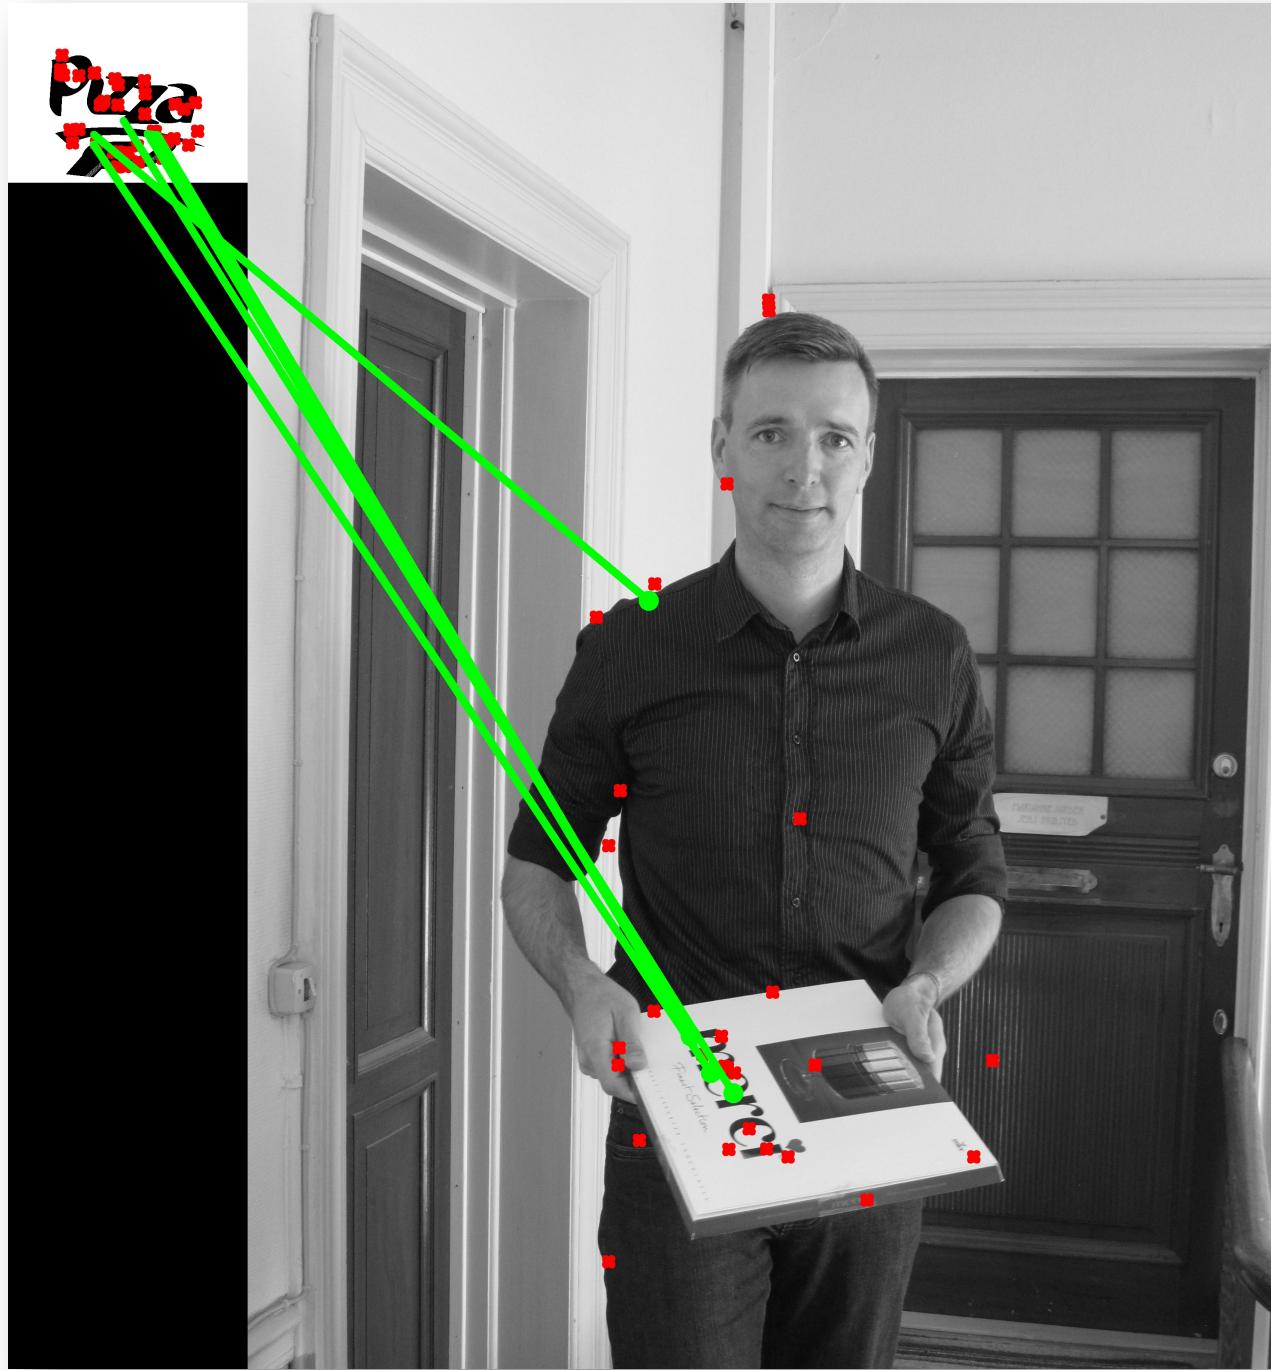
raw_matches = matcher.knnMatch(object_descriptors, scene_descriptors, k=2)
p1, p2, kp_pairs = filter_matches(object_keypoints, scene_keypoints, raw_matches)

MIN_MATCH_COUNT = 10

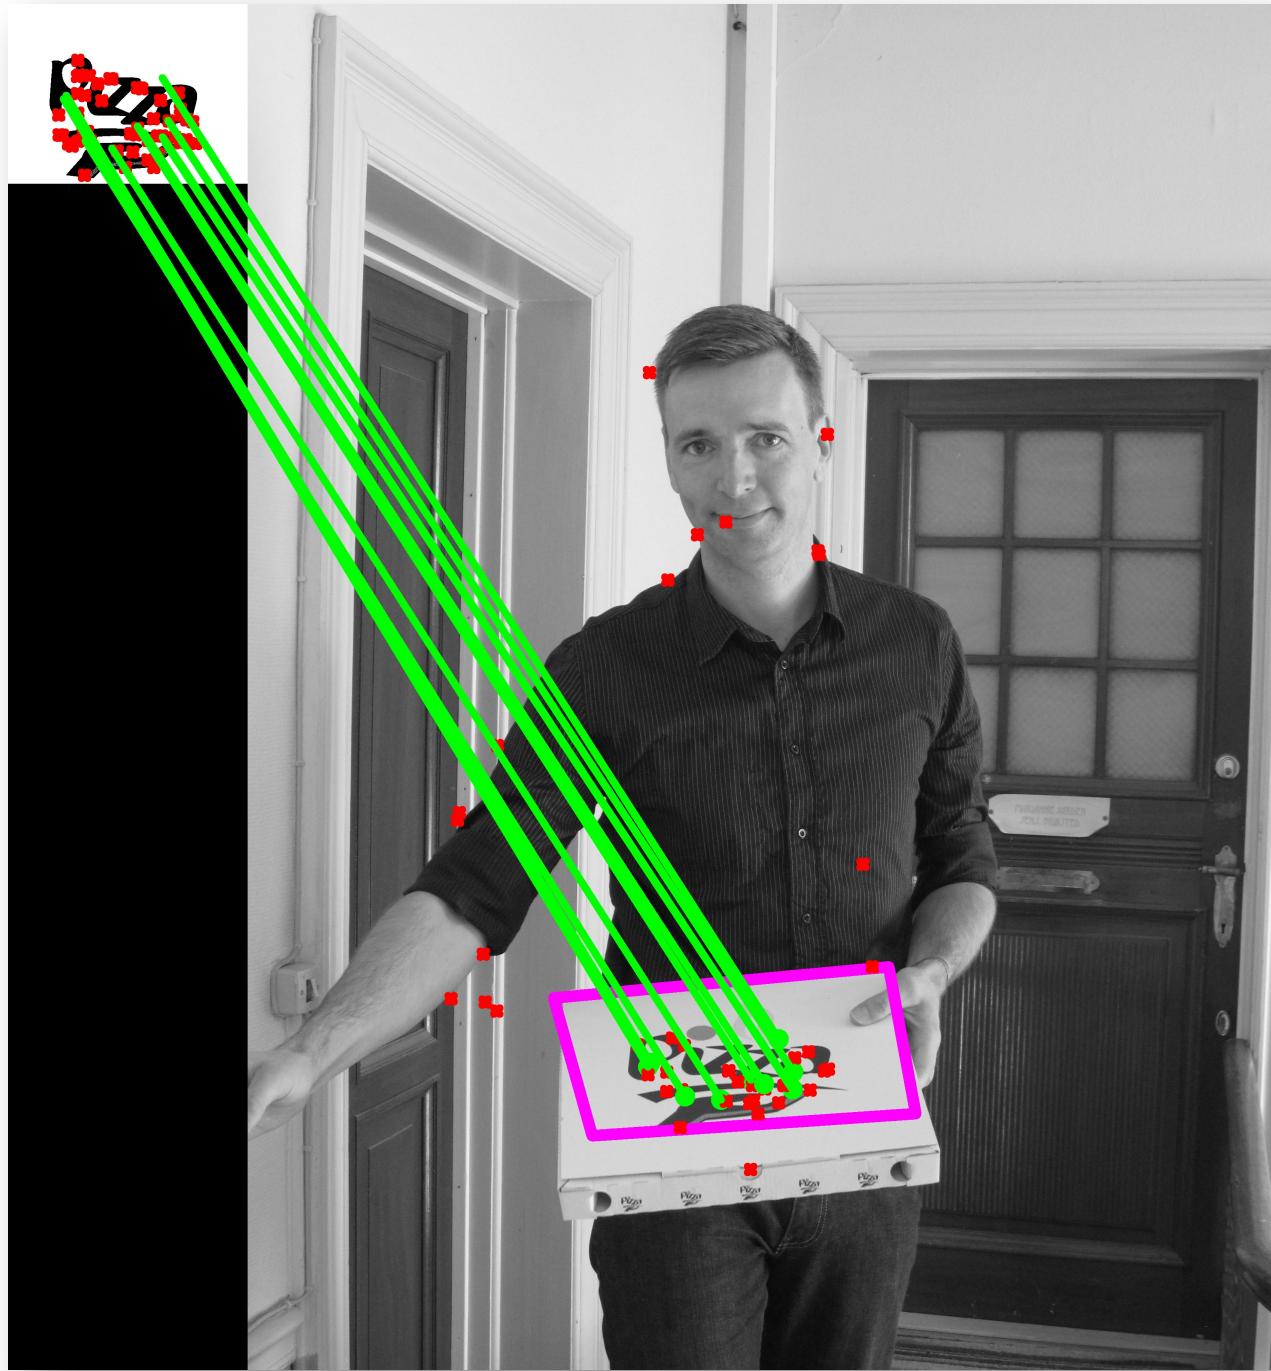
if len(kp_pairs)>MIN_MATCH_COUNT:
    H, status = CV2.findHomography(p1, p2, cv2.RANSAC, 5.0)
    match_lines_image = draw_matches(object_image, scene_image, kp_pairs, status, H)
    show('Match lines', match_lines_image)
```



Example code: [src/pizza.py](#)



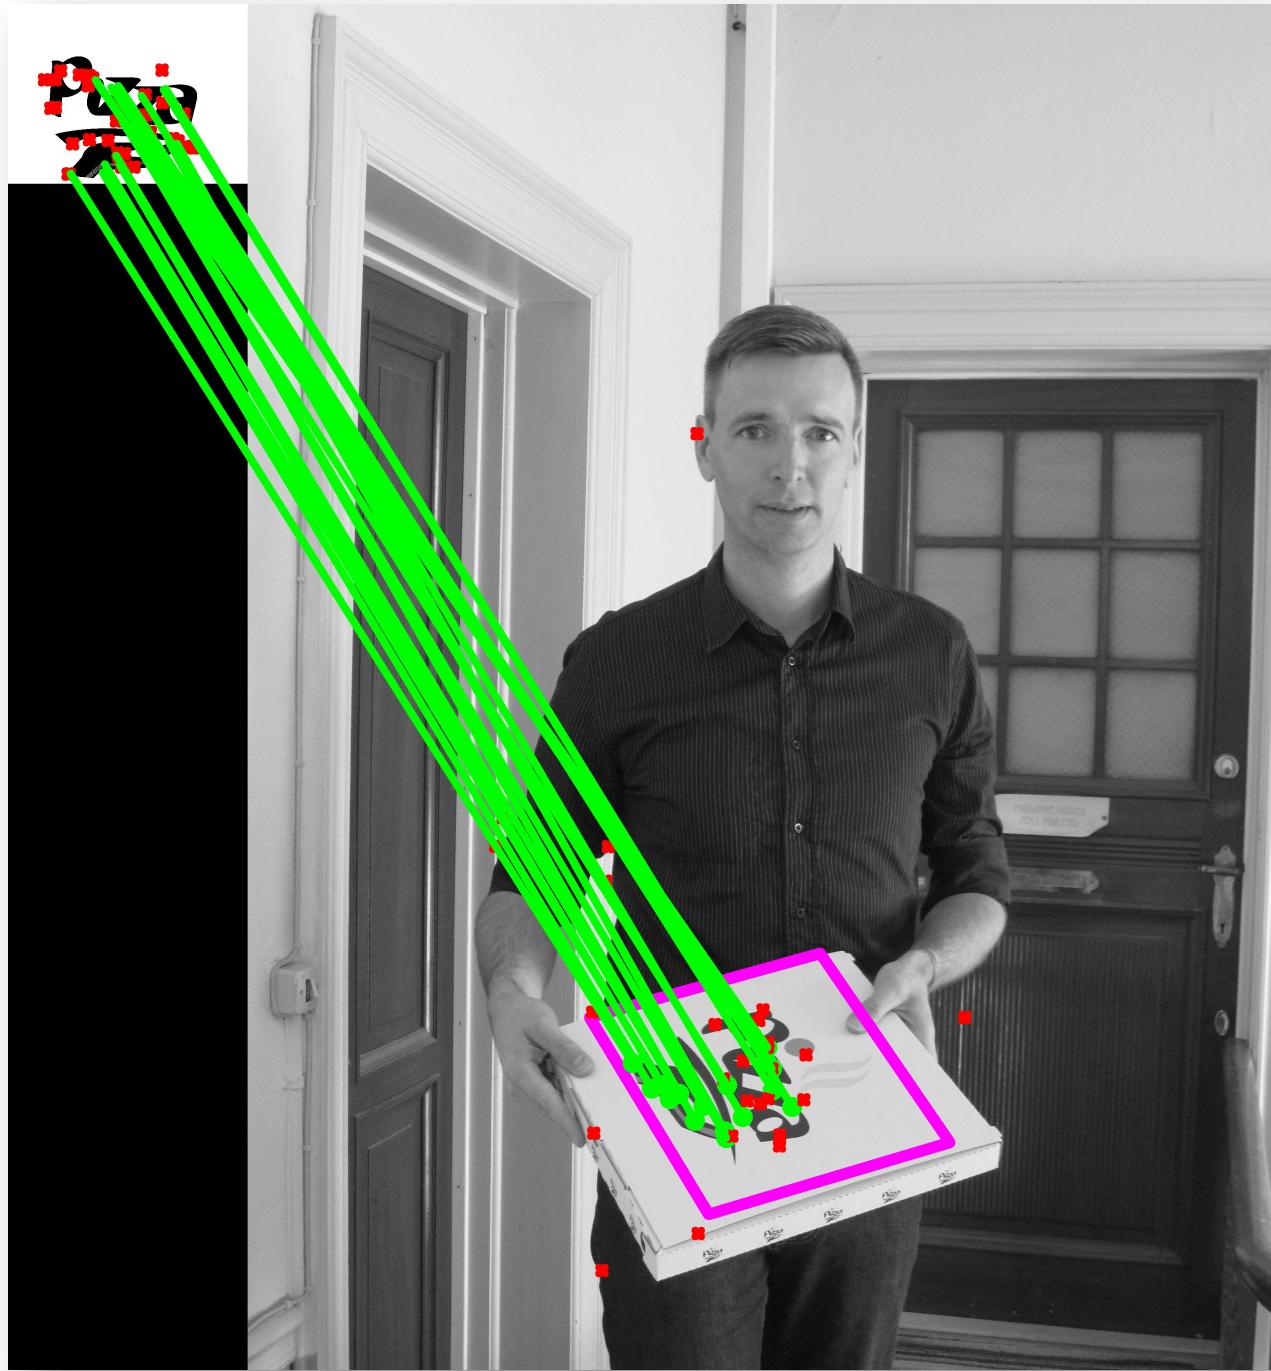
Example code: <src/pizza.py>



Example code: [src/pizza.py](#)



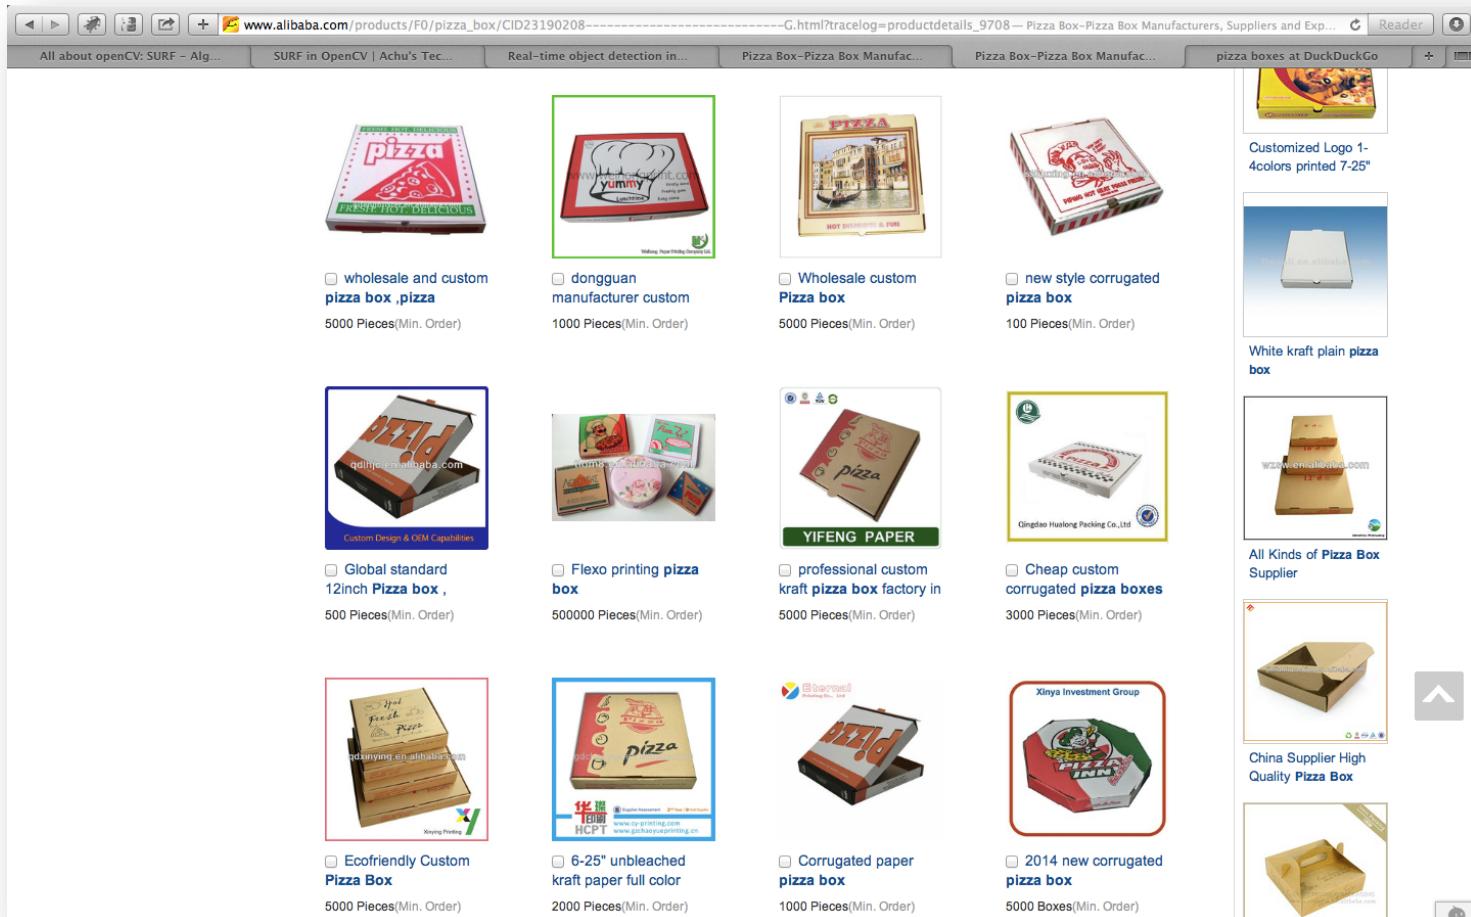
Example code: [src/pizza.py](#)



Example code: [src/pizza.py](#)

General Object Recognition

- Recognize the class of *all possible* pizza boxes



Visual Queries and Machine Learning

WINE

How Would You Do It?



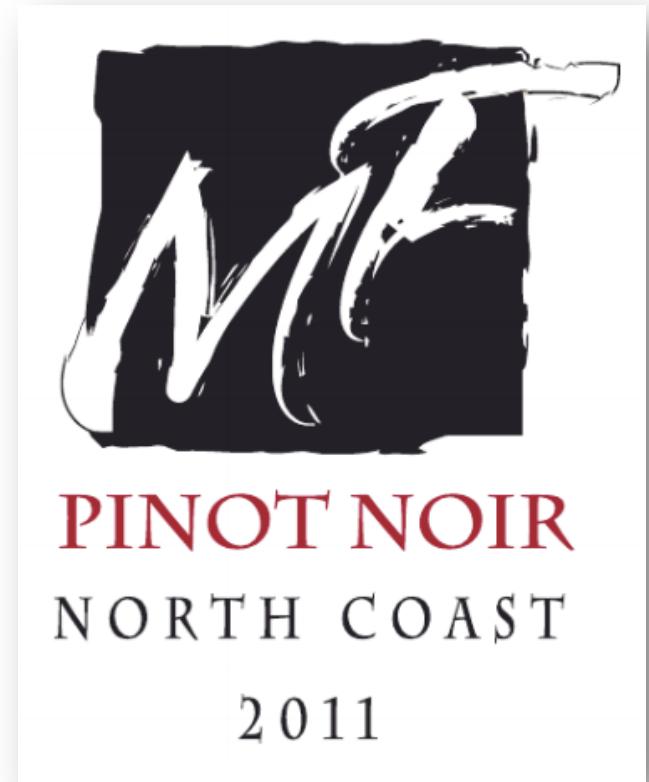
A screenshot of a mobile application interface for a wine database. At the top, there is a red header bar with icons for signal strength, battery level, and navigation. Below the header, the wine's name, 'The Four Graces', is displayed in large, elegant script. Underneath the name, the specific details are listed: 'Four Graces Dundee Hills Reserve Pinot Noir 2006'. Below this, the origin is noted as 'Willamette Valley, United States'. The average rating is shown as '4.0' with five stars, and the price is listed as '-' with a red wine icon. At the bottom of the screen, there is a red footer bar with five icons: a person, a bottle, a camera, a star, and a location pin.

Complexities



Photo

Partial
OCR
“likely”
*curved
lines*



*Official Logo, MF: Matthew Fritz
From <http://www.mfwine.com>*

Complexities

Wavy text – or pattern?

Distinct shape



OCR'able?

Artistic border
- or text?

How Would You Do It?

Example filters (feature detectors)

(Partial) OCR, e.g.

Letter frequencies (language)

Letter n-grams

Words (may be misread/missing)



"2011"

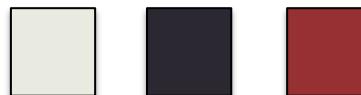
(2,0,1,1)
(2,0,1) (0,1,1)
(2,0) (0,1) (1,1)
(2) (0) (1) (1)

"2004"

(2,0,0,4)
(2,0,0) (0,0,4)
(2,0) (0,0) (0,4)
(2) (0) (0) (4)

Global features

e.g. Mean/dominant colours



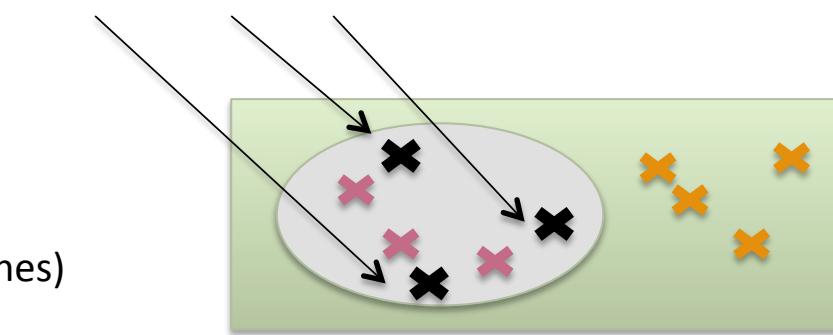
Best match against

Local features (Feature matching)

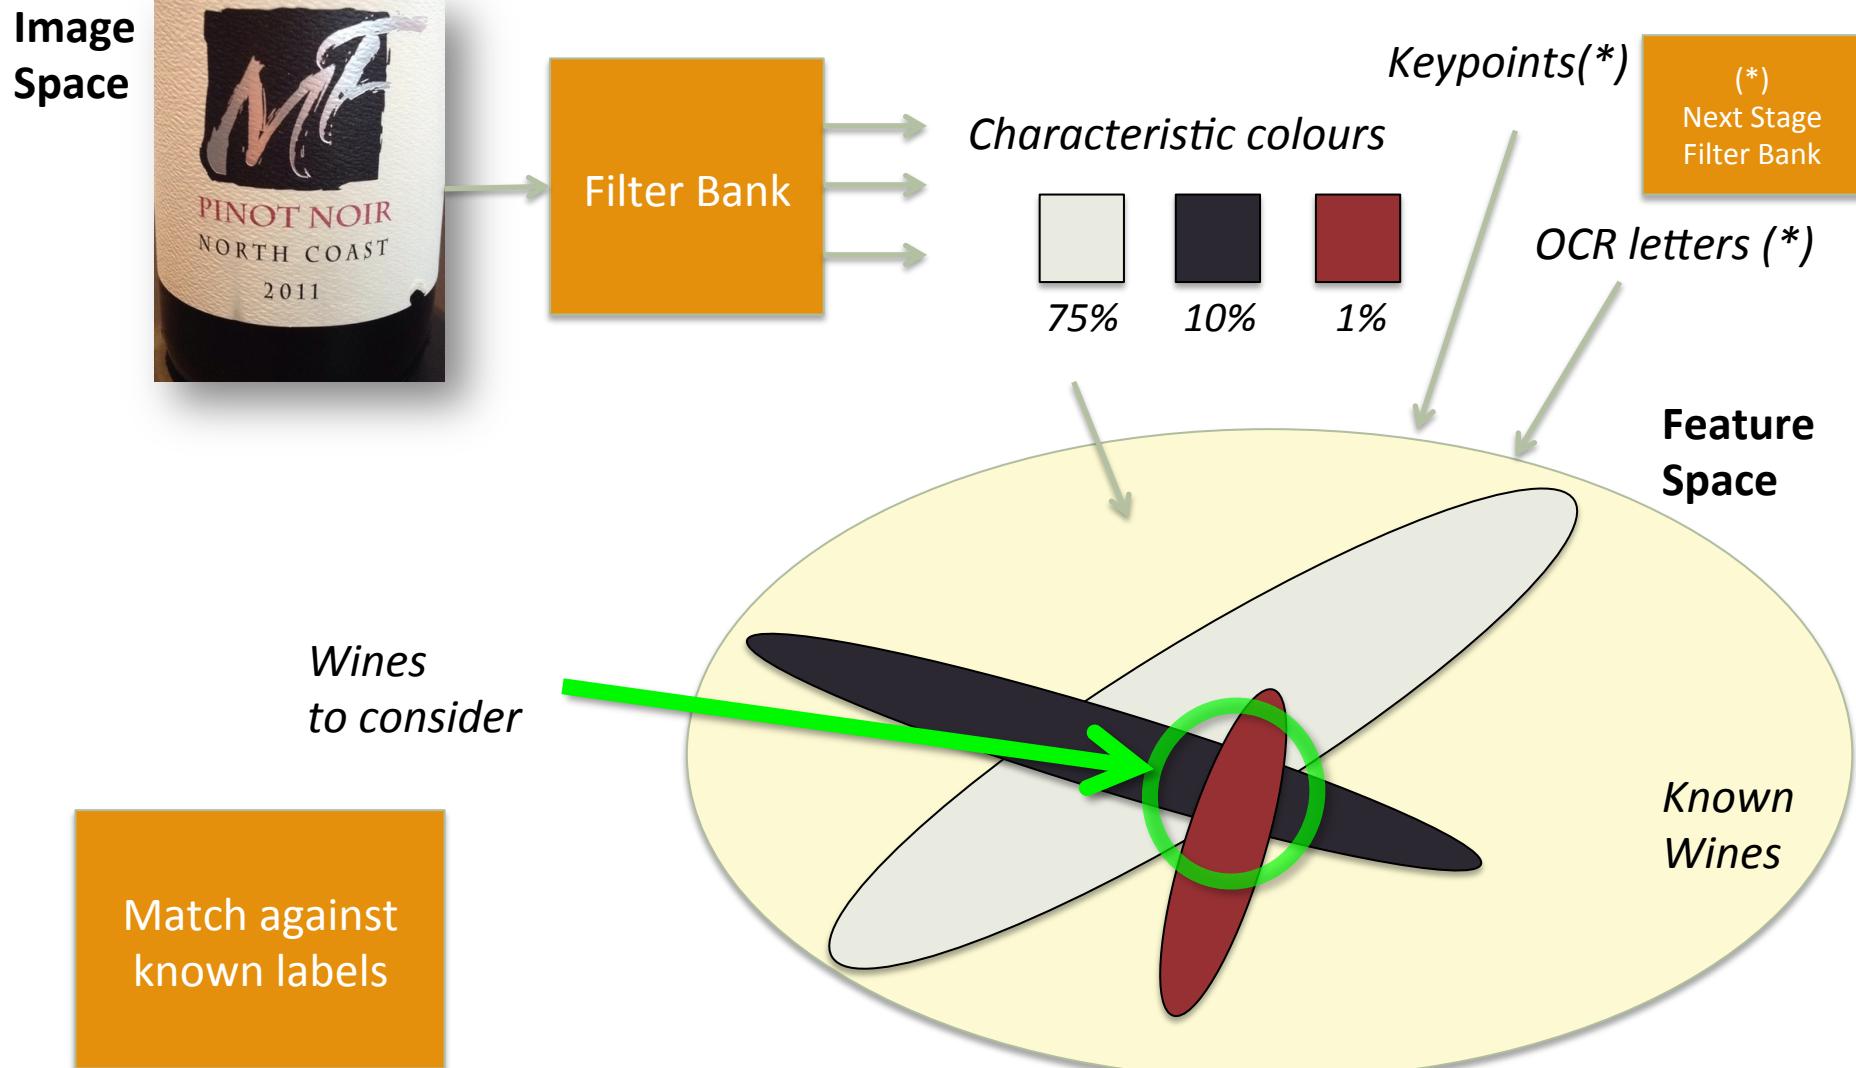
Classifier (trained from known images)

Database of producer names (known wines)

Match - nearest neighbour



Wine Matching Pipeline



Big Data in Your Pocket

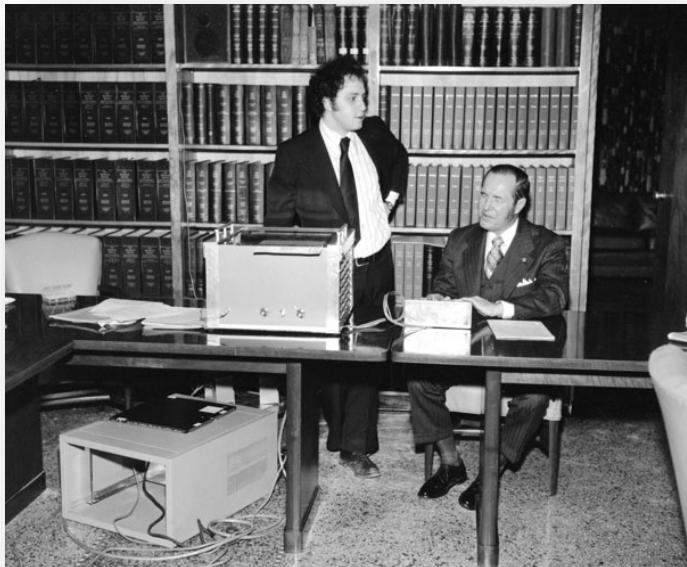
Kurzweil Reading Machine (1976)

64 KB RAM

¼ MHz CPU

Flatbed CCD Scanner

Software: Omnifont OCR, text-to-speech



January 1977 – Kurzweil unveils the Reading Machine at Iowa Commission for the Blind

<http://www.iowablindhistory.org/blindhistory/tools-and-technology-optical-character-recognition-scanners>

Samsung Galaxy S5 (2014)

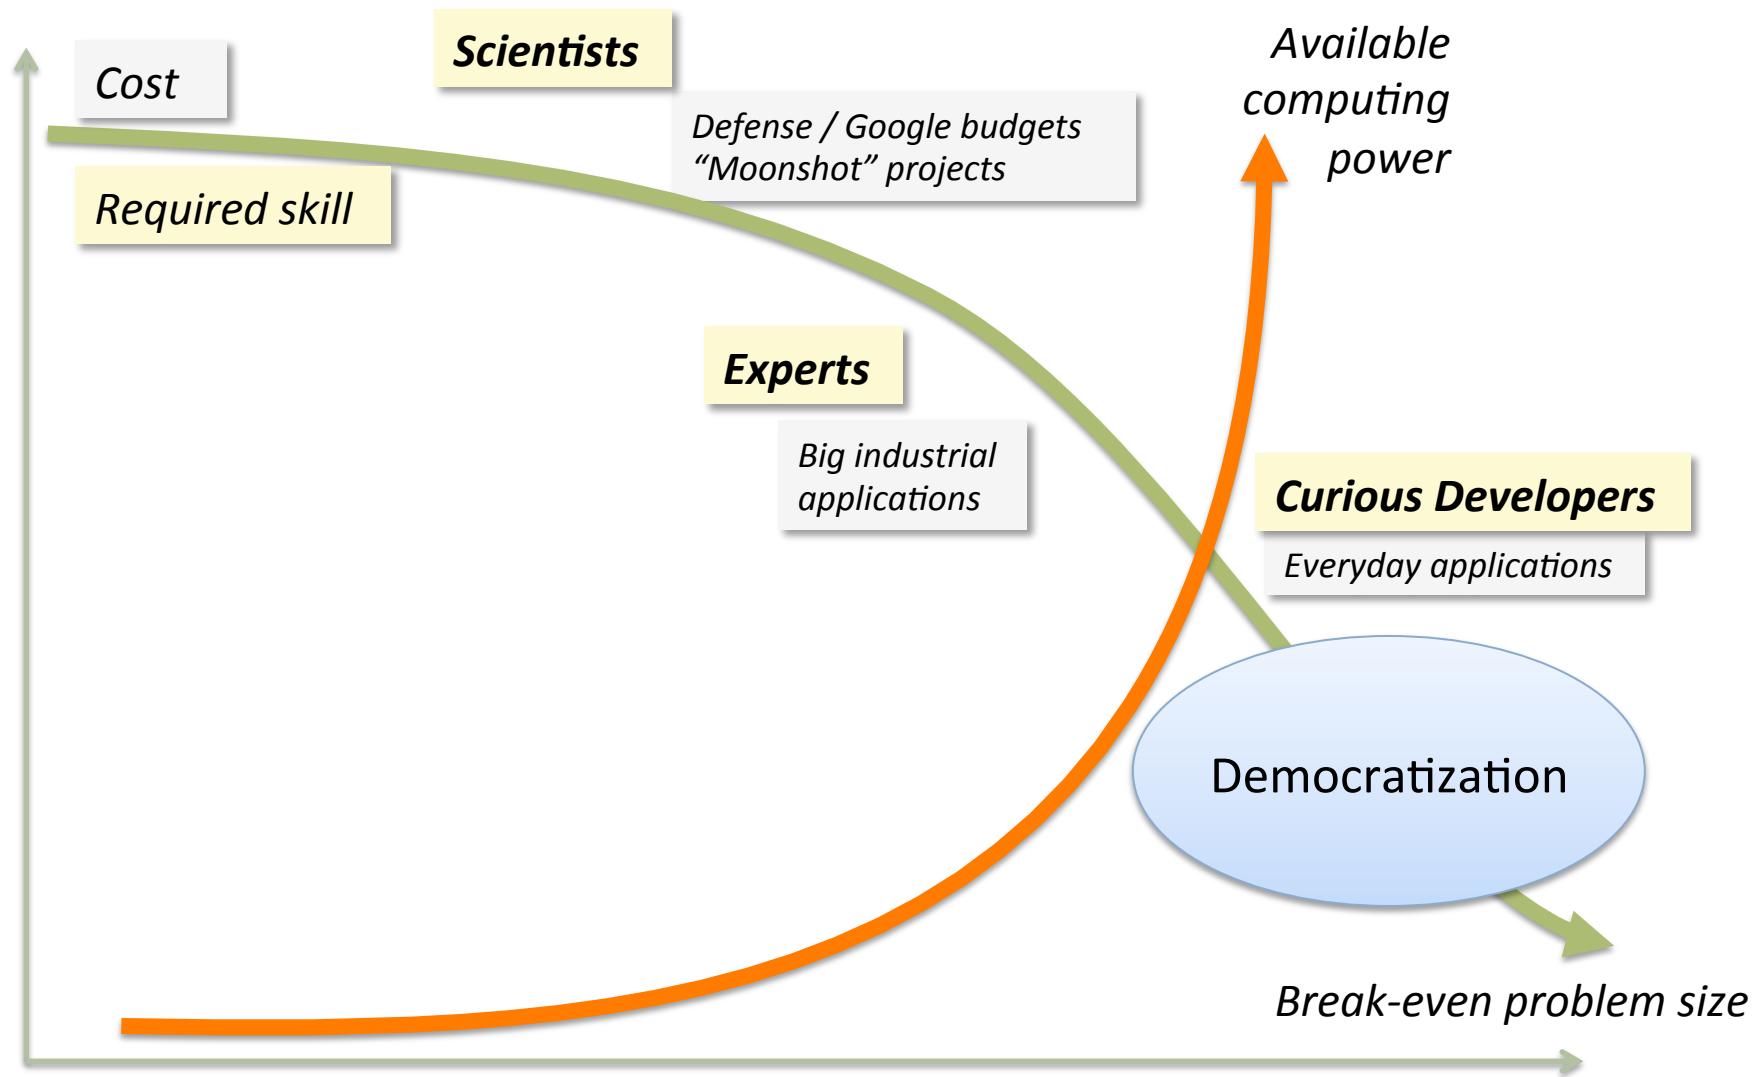
2 GB RAM

2.5 GHz quad-core CPU

16 megapixel camera



Computer Vision

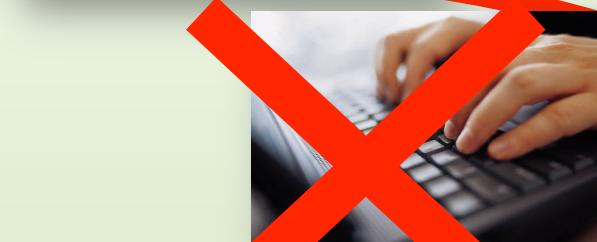


Re-imagine Everything

*Unstructured
Data*

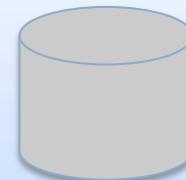


Human Data Processing



Four Graces
Dundas Hills Reserve
Pinot Noir
2006

*Structured
Data*



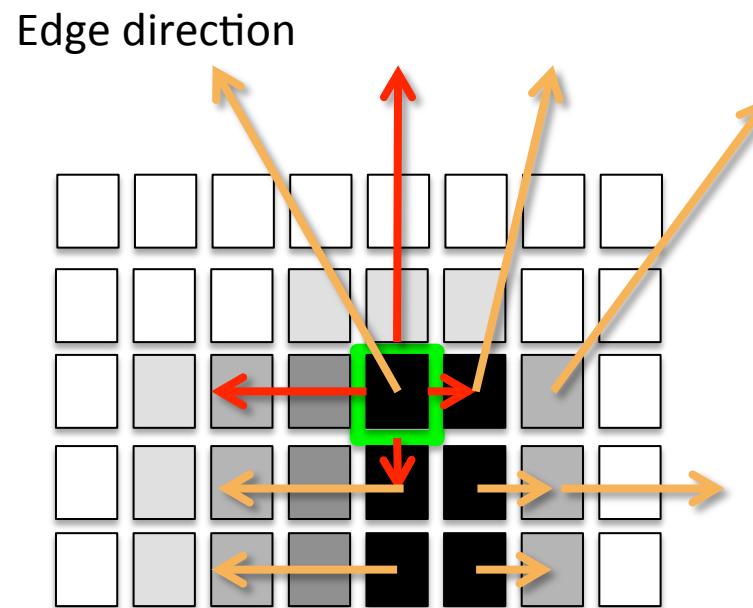
Machine Data Processing

Questions?

- Get in touch!
 - Martin Jul / martin@mjul.com / @mjul
- Get the slides and the code
 - <https://github.com/mjul/making-the-computer-see-ndc-2014>

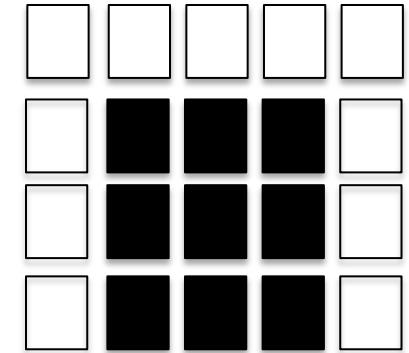
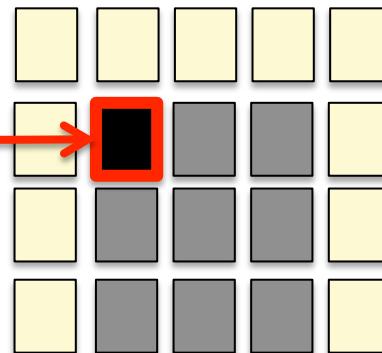
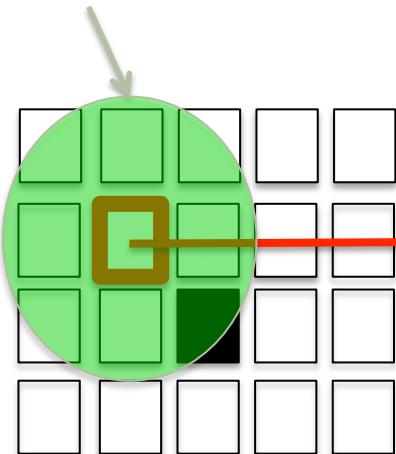


Canny Edge Detection



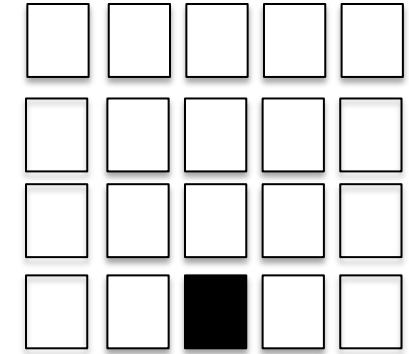
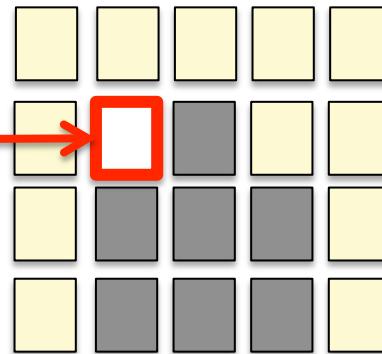
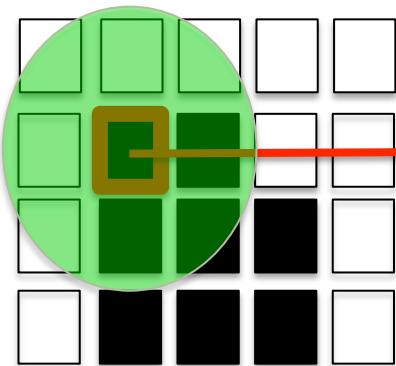
Dilation and Erosion

kernel



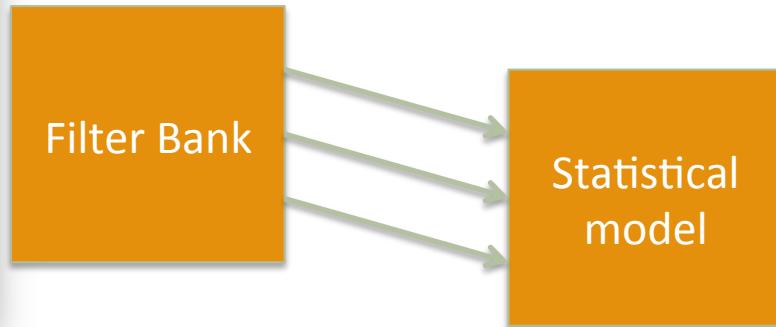
Erosion – \min over kernel neighbourhood

(8-bit grayscale: white = 255, black = 0)

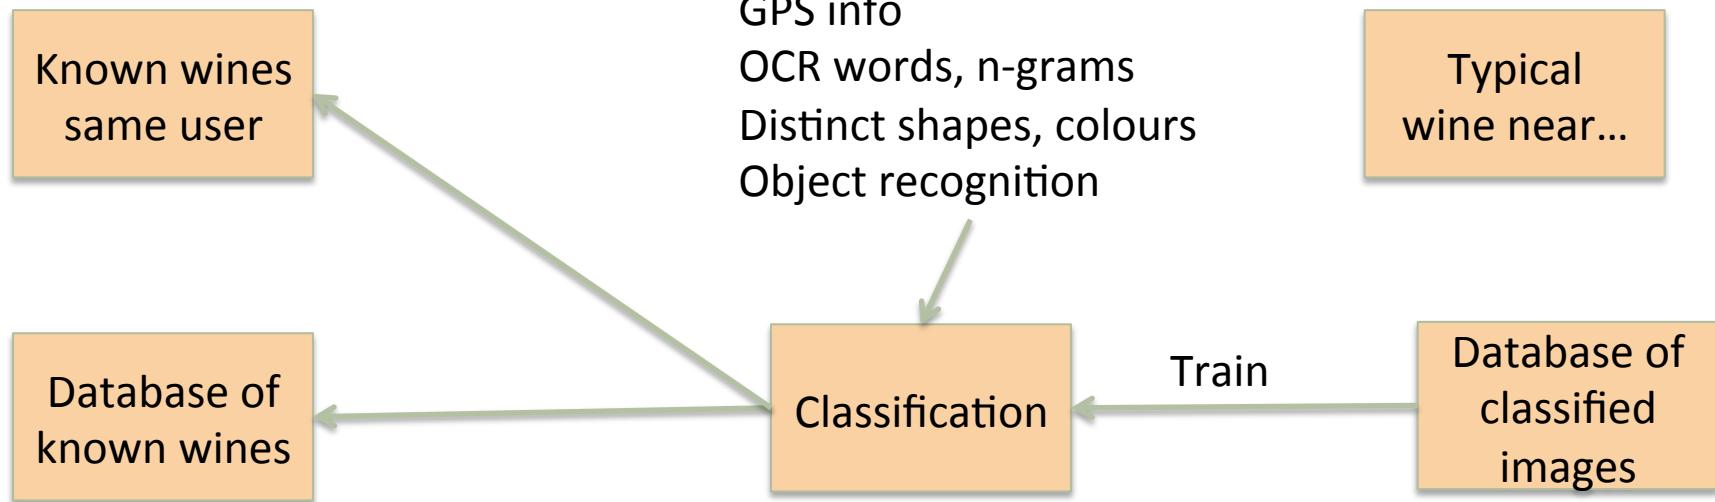


Dilation – \max over kernel neighbourhood

Machine Learning Wines



Query
Best match for
these features...



Machine Learning

